

MODULE - 5

4.1 Managing State

The Problem of State in Web Applications

Unified single process that is the typical desktop application, a web application consists of a series of disconnected HTTP requests to a web server where each request for a server page is essentially a request to run a separate program.




FIGURE 13.1 Desktop applications versus web applications

The web server sees only requests. The HTTP protocol does not, without programming intervention, distinguish two requests by one source from two requests from two different sources

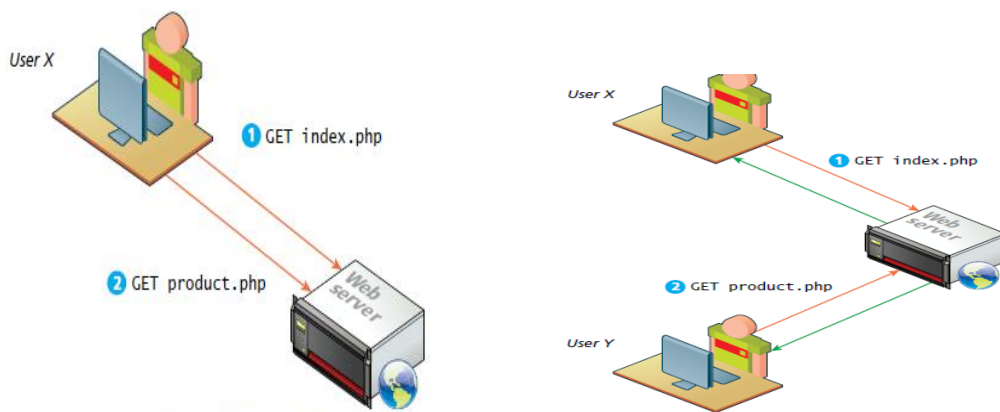


FIGURE 13.2 What the web server sees

... is for the server not really any different than ...

While the HTTP protocol disconnects the user's identity from his or her requests, there are many occasions when we want the web server to connect requests together. Consider the scenario of a web shopping cart, as shown in Figure 13.3. In such a case, the user (and the website owner) most certainly wants the server to recognize that the request to add an item to the cart and the subsequent request to check out and pay for the item in the cart are connected to the same individual.




FIGURE 13.3 What the user wants the server to see

In HTTP, we can pass information using:

- Query strings
- Cookies

4.2 Passing Information via Query Strings

A web page can pass query string information from the browser to the server using one of the two methods: a query string within the URL (GET) and a query string within the HTTP header (POST). Figure 13.4 reviews these two different approaches.

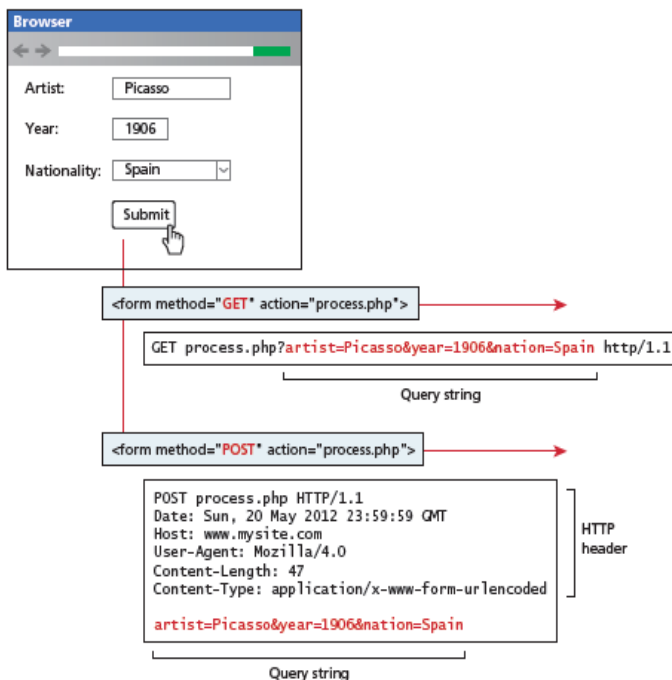


FIGURE 13.4 Recap of GET versus POST

4.3 Passing Information via the URL Path

While query strings are a vital way to pass information from one page to another, they do have a drawback. The URLs that result can be long and complicated. While for many users this is not that important, many feel that for one particular type of user, query strings are not ideal.

While there is some dispute about whether dynamic URLs (i.e., ones with query string parameters) or static URLs are better from a search engine result optimization (or SEO for search engine optimization) perspective, the consensus is that static URLs do provide some benefits with search engine result rankings. Many factors affect a page's ranking in a search engine, but the appearance of search terms within the URL does seem to improve its relative position.

Another benefit to static URLs is that users tend to prefer them. As we have seen, dynamic URLs (i.e., query string parameters) are a pretty essential part of web application development. How can we do without them? The answer is to rewrite the dynamic URL into a static one (and vice versa). This process is commonly called **URL rewriting**.

We can try doing our own rewriting. Let us begin with the following URL with its query string information:

`www.somedomain.com/DisplayArtist.php?artist=16`

One typical alternate approach would be to rewrite the URL to:

`www.somedomain.com/artists/16.php`

Notice that the query string name and value have been turned into path names.

One could improve this to make it more SEO friendly using the following:

`www.somedomain.com/artists/Mary-Cassatt`

The `mod_rewrite` module uses a rule-based rewriting engine that utilizes Perl compatible regular expressions to change the URLs so that the requested URL can be mapped or redirected to another URL internally.

4.4 Cookies

Cookies are a client-side approach for persisting state information. They are name=value pairs that are saved within one or more text files that are managed by the browser. These pairs accompany both server requests and responses within the HTTP header. While cookies cannot contain viruses, third-party tracking cookies have been a source of concern for privacy advocates.

Cookies were intended to be a long-term state mechanism. They provide website authors with a mechanism for persisting user-related information that can be stored on the user's computer and be managed by the user's browser. Cookies are not associated with a specific page but with the page's domain, so the browser and server will exchange cookie information no matter what page the user requests from the site. The browser manages the cookies for the different domains so that one domain's cookies are not transported to a different domain.

While cookies can be used for any state-related purpose, they are principally used as a way of maintaining continuity over time in a web application. One typical use of cookies in a website is to "remember" the visitor, so that the server can customize the site for the user. Some sites will use cookies as part of their shopping cart implementation so that items added to the cart will remain there even if the user leaves the site and then comes back later. Cookies are also frequently used to keep track of whether a user has logged into a site.

How Do Cookies Work?

While cookie information is stored and retrieved by the browser, the information in a cookie travels within the HTTP header. Figure 13.6 illustrates how cookies work. There are limitations to the amount of information that can be stored in a cookie (around 4K) and to the number of cookies for a domain (for instance, Internet Explorer 6 limited a domain to 20 cookies). HTTP cookies can also expire. That is, the browser will delete cookies that are beyond their expiry date (which is a configurable property of a cookie). If a cookie does not have an expiry date specified, the browser will delete it when the browser closes (or the next time it accesses the site). For this reason, some commentators will say that there are two types of cookies: session cookies and persistent cookies. A **session cookie** has no expiry stated and thus will be deleted at the end of the user browsing session. **Persistent cookies** have an expiry date specified; they will persist in the browser's cookie file until the expiry date occurs, after which they are deleted.

The most important limitation of cookies is that the browser may be configured to refuse them. As a consequence, sites that use cookies should not depend on their availability for critical features. Similarly, the user can also delete cookies or even tamper with the cookies, which may lead to some serious problems if not handled.

Several years ago, there was an instructive case of a website selling stereos and televisions that used a cookie-based shopping cart. The site placed not only the product identifier but also the product price in the cart. Unfortunately, the site then used the price in the cookie in the checkout. Several curious shoppers edited the price in the cookie stored on their computers, and then purchased some big-screen televisions for only a few cents!

Using Cookies

Like any other web development technology, PHP provides mechanisms for writing and reading cookies. Cookies in PHP are *created* using the `setcookie()` function and are *retrieved* using the `$_COOKIE` superglobal associative array. Below example illustrates the writing of a persistent cookie in PHP

```
<?php
// add 1 day to the current time for expiry time
$expiryTime = time()+60*60*24;
// create a persistent cookie
$name = "Username";
$value = "Ricardo";
setcookie($name, $value, $expiryTime);
?>
```

The `setcookie()` function also supports several more parameters, which further customize the new cookie. You can examine the online official PHP documentation for more information. The below example illustrates the reading of cookie values. Notice that when we read a cookie, we must also check to ensure that the cookie exists. In PHP, if the cookie has expired (or never existed in the first place), then the client's browser would not send anything, and so the `$_COOKIE` array would be blank.

```
<?php
if( !isset($_COOKIE['Username']) ) {
//no valid cookie found
}
```

```

else {
    echo "The username retrieved from the cookie is:";
    echo $_COOKIE['Username'];
}
?>

```

Persistent Cookie Best Practices

Many sites provide a “Remember Me” checkbox on login forms, which relies on the use of a persistent cookie. This login cookie would contain the user’s username but not the password. Instead, the login cookie would contain a random token; this random token would be stored along with the username in the site’s back-end database. Every time the user logs in, a new token would be generated and stored in the database and cookie.

Another common, nonessential use of cookies would be to use them to store user preferences. For instance, some sites allow the user to choose their preferred site color scheme or their country of origin or site language. In these cases, saving the user’s preferences in a cookie will make for a more contented user, but if the user’s browser does not accept cookies, then the site will still work just fine; at worst the user will simply have to reselect his or her preferences again.

Another common use of cookies is to track a user’s browsing behavior on a site. Some sites will store a pointer to the last requested page in a cookie; this information can be used by the site administrator as an analytic tool to help understand how users navigate through the site.

4.5 Serialization

Serialization is the process of taking a complicated object and reducing it down to zeros and ones for either storage or transmission. Later that sequence of zeros and ones can be reconstituted into the original object.

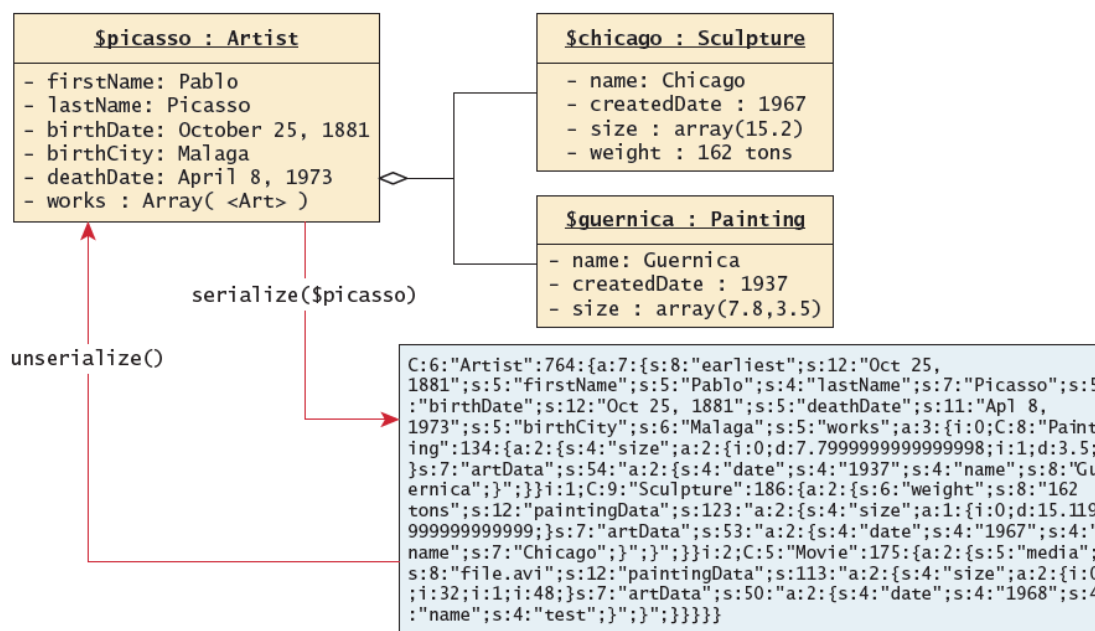


FIGURE 13.7 Serialization and deserialization

In PHP objects can easily be reduced down to a binary string using the `serialize()` function. The resulting string is a binary representation of the object and therefore may contain unprintable characters. The string can be reconstituted back into an object using the `unserialize()` method. While arrays, strings, and other primitive types will be serializable by default, classes of our own creation must implement the `Serializable` interface shown below. Which requires adding implementations for `serialize()` and `unserialize()` to any class that implements this interface.

```

interface Serializable {
    /* Methods */
    public function serialize();
    public function unserialize($serialized);
}

```

The example shows how the `Artist` class must be modified to implement the `Serializable` interface by adding the `implements` keyword to the class definition and adding implementations for the two methods.

listing 13.4 Artist class modified to implement the `Serializable` interface

```

class Artist implements Serializable {
    //...
    // Implement the Serializable interface methods
}

```

```

public function serialize() {
    // use the built-in PHP serialize function
    return serialize(
        array("earliest" => self::$earliestDate,
            "first" => $this->firstName,
            "last" => $this->lastName,
            "bdate" => $this->birthDate,
            "ddate" => $this->deathDate,
            "bcity" => $this->birthCity,
            "works" => $this->artworks
        );
    );
}

public function unserialize($data) {
    // use the built-in PHP unserialize function
    $data = unserialize($data);
    self::$earliestDate = $data['earliest'];
    $this->firstName = $data['first'];
    $this->lastName = $data['last'];
    $this->birthDate = $data['bdate'];
    $this->deathDate = $data['ddate'];
    $this->birthCity = $data['bcity'];
    $this->artworks = $data['works'];
}
//...
}

```

If the data above is assigned to `$data`, then the following line will instantiate a new object identical to the original:

```
$picassoClone = unserialize($data);
```

Application of Serialization

Since each request from the user requires objects to be reconstituted, using serialization to store and retrieve objects can be a rapid way to maintain state between requests. At the end of a request you store the state in a serialized form, and then the next request would begin by deserializing it to reestablish the previous state.

4.6 Session State

Session state is a server-based state mechanism that lets web applications store and retrieve objects of any type for each unique user session. That is, each browser session has its own session state stored as a serialized file on the server, which is deserialized and loaded into memory as needed for each request




FIGURE 13.8 Session state

Because server storage is a finite resource, objects loaded into memory are released when the request completes, making room for other requests and their session objects. This means there can be more active sessions on disk than in memory at any one time. Session state is ideal for storing more complex (but not too complex . . . more on that later) objects or data structures that are associated with a user session. The classic example is a shopping cart. While shopping carts could be implemented via cookies or query string parameters, it would be quite complex and cumbersome to do so.

It can be accessed via the `$_SESSION` variable, but unlike the other superglobals, you have to take additional steps in your own code in order to use the `$_SESSION` superglobal. To use sessions in a script, you must call the `session_start()` function at the beginning of the script as shown in Listing 13.5.

```
<?php
session_start();
if ( isset($_SESSION['user']) ) {
    // User is logged in
}
else {
    // No one is logged in (guest)
}
?>
```

In this example, we differentiate a logged-in user from a guest by checking for the existence of the `$_SESSION['user']` variable. Session state is typically used for storing information that needs to be preserved across multiple requests by the same user. Since each user session has its own session state collection, it should not be used to store large amounts of information because this will consume very large amounts of server memory as the number of active sessions increase. As well, since session information does eventually time out, one should always check if an item retrieved from session state still exists before using the retrieved object. If the session object does not yet exist (either because it is the first time the user has requested it or because the session has timed out), one might generate an error, redirect to another page, or create the required object using the lazy initialization approach as shown in Listing 13.6. In this example `ShoppingCart` is a user defined class. Since PHP sessions are serialized into files, one must ensure that any

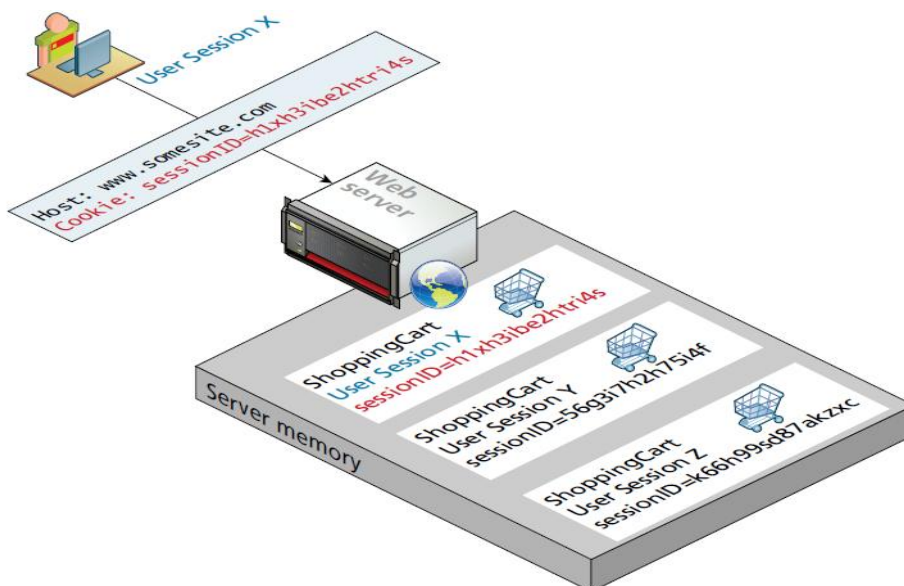
listing 13.5 Accessing session state

```
<?php
include_once("ShoppingCart.class.php");
session_start();
// always check for existence of session object before accessing it
if ( !isset($_SESSION["Cart"]) ) {
    //session variables can be strings, arrays, or objects, but
    // smaller is better
    $_SESSION["Cart"] = new ShoppingCart();
}
$cart = $_SESSION["Cart"];
?>
```

How Does Session State Work?

The first thing to know about session state is that it works within the same HTTP context as any web request. The server needs to be able to identify a given HTTP request with a specific user request. Since HTTP is stateless, some type of user/session identification system is needed.

In PHP, this is a unique 32-byte string that is by default transmitted back and forth between the user and the server via a session cookie as shown in Figure 13.9.



For a brand new session, PHP assigns an initially empty dictionary-style collection that can be used to hold any state values for this session. When the request processing is finished, the session state is saved to some type of state storage mechanism, called a session state provider (discussed in next section). Finally, when a new request is received for an already existing session; the session's dictionary collection is filled with the previously saved session data from the session state provider.

Session Storage and Configuration

It is possible to configure many aspects of sessions including where the session files are saved. For a complete listing refer to the session configuration options in **php.ini**. The decision to save sessions to files rather than in memory (like ASP.NET) addresses the issue of memory usage that can occur on shared hosts as well as persistence between restarts. Many sites run in commercial hosting environments that are also hosting many other sites. For instance, one of the book author's personal sites (**randyconnolly.com**, which is hosted by **discountasp.net**) is, according to a Reverse IP Domain Check, on a server that was hosting 68 other sites when this chapter was being written. Inexpensive web hosts may sometimes stuff hundreds or even thousands of sites on each machine. In such an environment, the server memory that is allotted per web application will be quite limited. And remember that for each application, server memory may be storing not only session information, but pages being executed, and caching information, as shown in Figure.

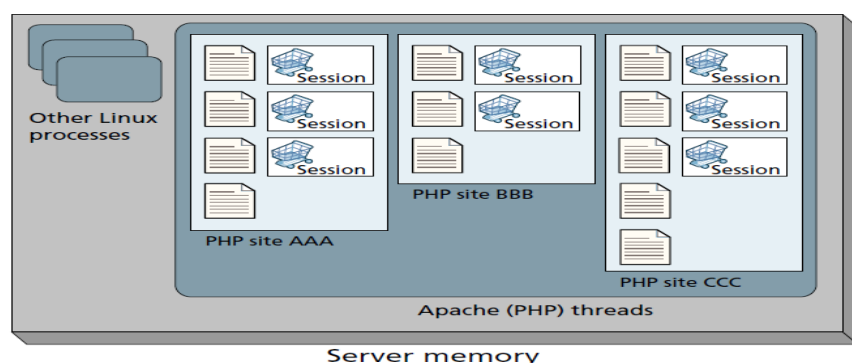


FIGURE 13.10 Applications and server memory

On a busy server hosting multiple sites, it is not uncommon for the Apache application process to be restarted on occasion. If the sessions were stored in memory, the sessions would all expire, but as they are stored into files, they can be instantly recovered as though nothing happened. This can be an issue in environments where sessions are stored in memory (like ASP.NET), or a custom session handler is involved. One downside to storing the sessions in files is degradation in performance compared to memory storage, but the advantages, it was decided, outweigh those challenges.

Higher-volume web applications often run in an environment in which multiple web servers (also called a web farm) are servicing requests. Each incoming request is forwarded by a load balancer to any one of the available servers in the farm. In such a situation the in-process session state will not work, since one server may service one request for a particular session, and then a completely different server may service the next request for that session, as shown in Figure 13.11.




FIGURE 13.11 Web farm

There are a number of different ways of managing session state in such a web farm situation, some of which can be purchased from third parties. There are effectively two categories of solution to this problem.

1. Configure the load balancer to be "session aware" and relate all requests using a session to the same server.
2. Use a shared location to store sessions, either in a database, memcache, or some other shared session state mechanism as seen in Figure 13.12.




FIGURE 13.12 Shared session provider

Using a database to store sessions is something that can be done programmatically, but requires a rethinking of how sessions are used. Code that was written to work on a single server will have to be changed to work with sessions in a shared database, and therefore is cumbersome. The other alternative is to configure PHP to use memcache on a shared server (covered in Section 13.8). To do this you must have PHP compiled with memcache enabled; if not, you may need to install the module. Once installed, you must change the `php.ini` on all servers to utilize a shared location, rather than local files as shown in Listing 13.7.

[Session]

; Handler used to store/retrieve data.

```
session.save_handler = memcache
```

```
session.save_path = "tcp://sessionServer:11211"
```

4.7 JavaScript Pseudo-Classes

Using Object Literals

An array in JavaScript can be instantiated with elements in the following way:

```
var daysofWeek = ["sun","mon","tue","wed","thu","fri","sat"];
```

An object can be instantiated using the similar concept of **object literals**: that is, an object represented by the list of key-value pairs with colons between the key and value with commas separating key-value pairs. A dice object, with a string to hold the color and an array containing the values representing each side (face), could be defined all at once using object literals as follows:

```
var oneDie = { color : "FF0000", faces : [1,2,3,4,5,6] };
```

Once defined, these elements can be accessed using dot notation. For instance, one could change the color to blue by writing:

```
oneDie.color="0000FF";
```

Emulate Classes through Functions

Although a formal *class* mechanism is not available to us in JavaScript, it is possible to get close by using functions to encapsulate variables and methods together,

```
function Die(col) {
  this.color=col;
  this.faces=[1,2,3,4,5,6];
}
```


The **this** keyword inside of a function refers to the instance, so that every reference to internal properties or methods manages its own variables, as is the case with PHP. One can create an instance of the object as follows, very similar to PHP.

```
var oneDie = new Die("0000FF");
```

Adding Methods to the Object

To define a method in an object's function one can either define it internally, or use a reference to a function defined outside the class. External definitions can quickly cause namespace conflict issues, since all method names must remain conflict free with all other methods for other classes. For this reason, one technique for adding a method inside of a class definition is by assigning an anonymous function to a variable.

```
function Die(col) {  
  this.color=col;  
  this.faces=[1,2,3,4,5,6];  
// define method randomRoll as an anonymous function  
  this.randomRoll = function() {  
    var randNum = Math.floor((Math.random() * this.faces.length)+ 1);  
    return faces[randNum-1];  
  };  
}  
var oneDie = new Die("0000FF");  
console.log(oneDie.randomRoll() + " was rolled");
```

Using Prototypes

Prototypes are an essential syntax mechanism in JavaScript, and are used to make JavaScript behave more like an object-oriented language. The prototype properties and methods are defined *once* for all instances of an *object*.

```
// Start Die Class  
function Die(col) {  
  this.color=col;  
  this.faces=[1,2,3,4,5,6];  
}  
Die.prototype.randomRoll = function() {  
  var randNum = Math.floor((Math.random() * this.faces.length) + 1);  
  return faces[randNum-1];  
};  
// End Die Class
```

More about Prototypes

Even experienced JavaScript programmers sometimes struggle with the prototype concept. It should be known that every object (and method) in JavaScript has a prototype.

A prototype is an object from which other objects inherit. The above definition sounds almost like a class in an object-oriented language, except that a prototype is itself an *object*, whereas in other oriented-oriented languages

A class is an abstraction, not an object. Despite this distinction, you can make use of a function's prototype object, and assign properties or methods to it that are then available to any new objects that are created.

The below example defines just such a method, named **countChars**, that takes a character as a parameter.

```
listing 15.4 Adding a method named countChars to the String class  
String.prototype.countChars = function (c) {  
  var count=0;  
  for (var i=0;i<this.length;i++) {  
    if (this.charAt(i) == c)  
      count++;  
  }  
  return count;  
}
```

Now any new instances of String will have this method available to them (created using the new keyword), while existing strings will not. You could use the new method on any strings instantiated after the prototype definition was added. For instance the following example will output Hello World has 3 letter l's.

```
var hel = "Hello World";

console.log(hel + "has" + hel.countChars("l") + " letter l's");
```

This technique is also useful to assign properties to a pseudo-class that you want available to all instances. Imagine an array of all valid characters attached to some custom string class. Again using prototype you could define such a list.

```
CustomString.prototype.validChars = ["A","B","C"];
```

4.8 JSON

JSON stands for JavaScript Object Notation; its use is not limited to JavaScript. It provides a more concise format than XML to represent data. It was originally designed to provide a lightweight serialization format to represent objects in JavaScript. While it doesn't have the validation and readability of XML, it has the advantage of generally requiring significantly fewer bytes to represent data than XML, which in the web context is quite significant.

Like XML, JSON is a data serialization format. That is, it is used to represent object data in a text format so that it can be transmitted from one computer to another. Many REST web services encode their returned data in the JSON data format instead of XML.

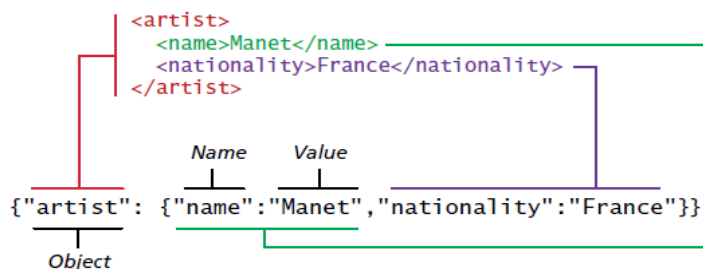


FIGURE 17.6 Sample JSON

Just like XML, JSON data can be nested to represent objects within objects. In general JSON data will have all white space removed to reduce the number of bytes traveling across the network.

```
{
  "paintings": [
    {
      "id": 290,
      "title": "Balcony",
      "artist": {
        "name": "Manet",
        "nationality": "France"
      },
      "year": 1868,
      "medium": "Oil on canvas"
    },
  ],
}
```

```

{
  "id":192,
  "title":"The Kiss",
  "artist":{
    "name":"Klimt",
    "nationality":"Austria"
  },
  "year":1907,
  "medium":"Oil and gold on canvas"
},
{
  "id":139,
  "title":"The Oath of the Horatii",
  "artist":{
    "name":"David",
    "nationality":"France"
  },
  "year":1784,
  "medium":"Oil on canvas"
}
]
}

```

Using JSON in JavaScript

Since the syntax of JSON is the same used for creating objects in JavaScript, it is easy to make use of the JSON format in JavaScript:

```

<script>
var a = {"artist": {"name":"Manet","nationality":"France"}};
alert(a.artist.name + " " + a.artist.nationality);
</script>

```

JSON information will be contained within a string, and the `JSON.parse()` function can be used to transform the string containing the JSON data into a JavaScript object:

```

var text = '{"artist": {"name":"Manet","nationality":"France"}}';
var a = JSON.parse(text);
alert(a.artist.nationality);

```

The jQuery library also provides a JSON parser that will work with all browsers (the `JSON.parse()` function is not available on older browsers):

```

var artist = jQuery.parseJSON(text);

```

JavaScript also provides a mechanism to translate a JavaScript object into a JSON string:

```

var text = JSON.stringify(artist);

```

Using JSON in PHP

PHP comes with a JSON extension and as of version 5.2 of PHP; the JSON extension is bundled and compiled into PHP by default. Converting a JSON string into a PHP object is quite straightforward:

```

<?php
// convert JSON string into PHP object
$text = '{"artist": {"name":"Manet","nationality":"France"}}';
$anObject = json_decode($text);
echo $anObject->artist->nationality;
// convert JSON string into PHP associative array
$anArray = json_decode($text, true);
echo $anArray['artist']['nationality'];
?>

```

`json_decode()` function can return either a PHP object or an associative array. Since JSON data is often coming from an external source, one should always check for parse errors before using it, which can be done via the `json_last_error()` function:

```

<?php
// convert JSON string into PHP object
$text = '{"artist": {"name": "Manet", "nationality": "France"}}';
$anObject = json_decode($text);
// check for parse errors
if (json_last_error() == JSON_ERROR_NONE) {
    echo $anObject->artist->nationality;
}
?>

```

To go the other direction (i.e., to convert a PHP object into a JSON string), you can use the `json_encode()` function.

```

// convert PHP object into a JSON string
$text = json_encode($anObject);

```

4.9 AJAX

Asynchronous JavaScript with XML (AJAX) is a term used to describe a paradigm that allows a web browser to send messages back to the server without interrupting the flow of what's being shown in the browser. This makes use of a browser's multi-threaded design and lets one thread handle the browser and interactions while other threads wait for responses to asynchronous requests.

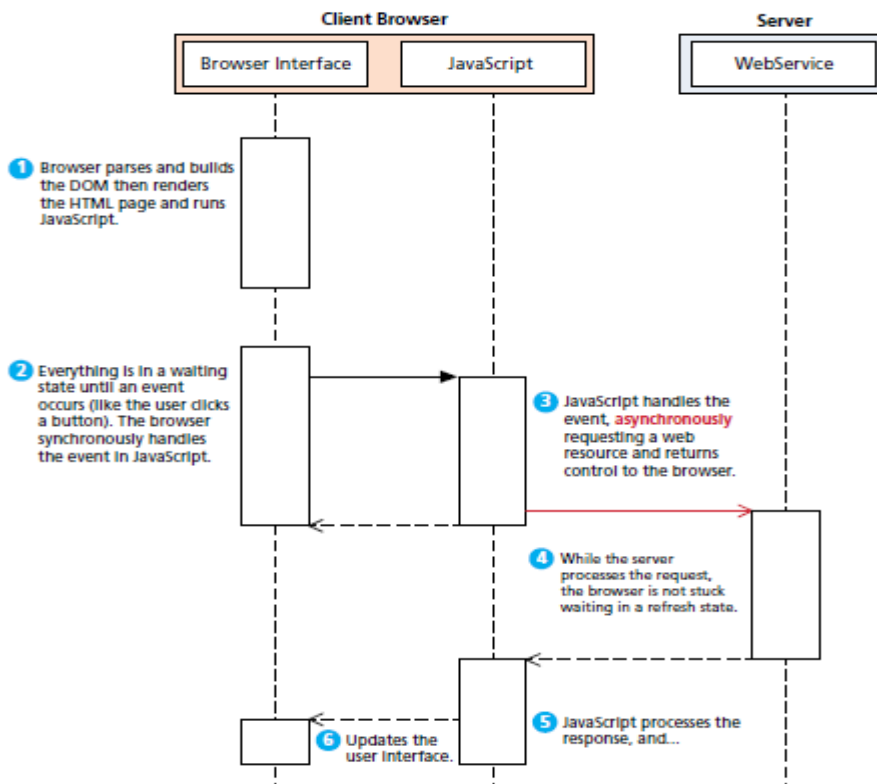


FIGURE 15.8 UML sequence diagram of an AJAX request

Responses to asynchronous requests are caught in JavaScript as events. The events can subsequently trigger changes in the user interface or make additional requests. This differs from the typical synchronous requests we have seen thus far, which require the entire web page to refresh in response to a request. Another way to contrast AJAX and synchronous JavaScript is to consider a webpage that displays the current server time as shown in the below figure.

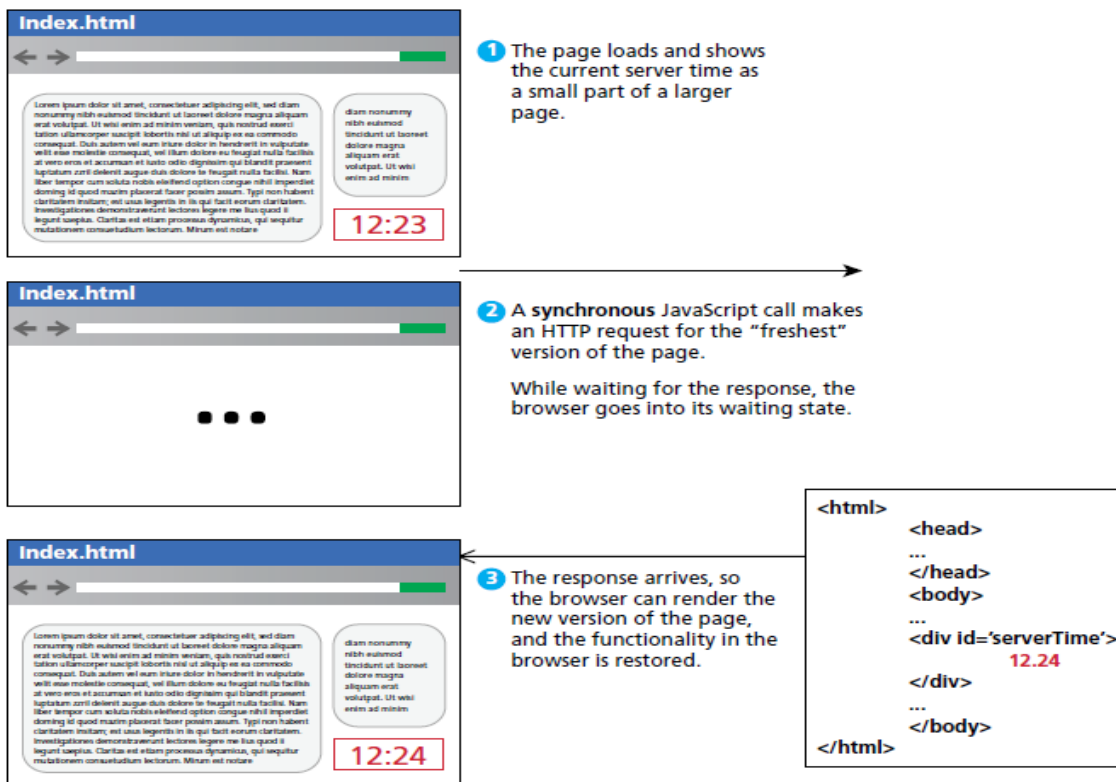


FIGURE 15.9 Illustration of a synchronous implementation of the server time web page.

If implemented synchronously, the entire page has to be refreshed from the server just to update the displayed time. During that refresh, the browser enters a waiting state, so the user experience is interrupted (yes, you could implement a refreshing time using pure JavaScript, but for illustrative purposes, imagine it's essential to see the server's time).

In contrast, consider the very simple asynchronous implementation of the server time, where an AJAX request updates the server time in the background as illustrated in Figure

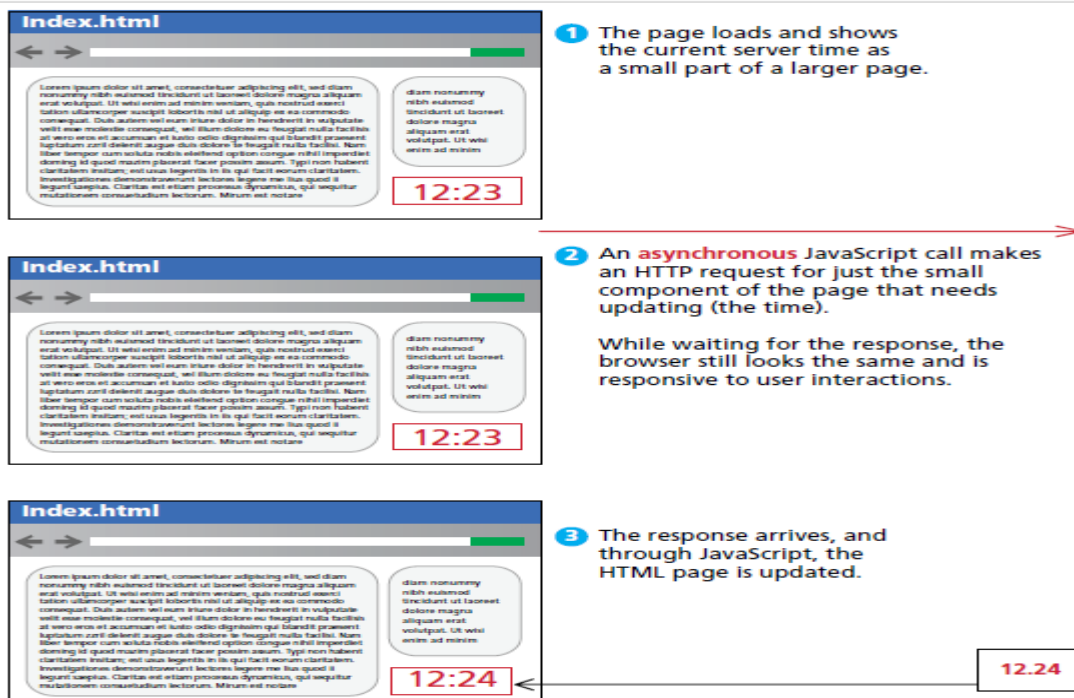


FIGURE 15.10 Illustration of an AJAX implementation of the server time widget

Making Asynchronous Requests

jQuery provides a family of methods to make asynchronous requests. Consider for instance the very simple server time page described above. If the URL `currentTime.php` returns a single string and you want to load that value asynchronously

into the `<div id="timeDiv">` element, you could write:

```
$("#timeDiv").load("currentTime.php");
```

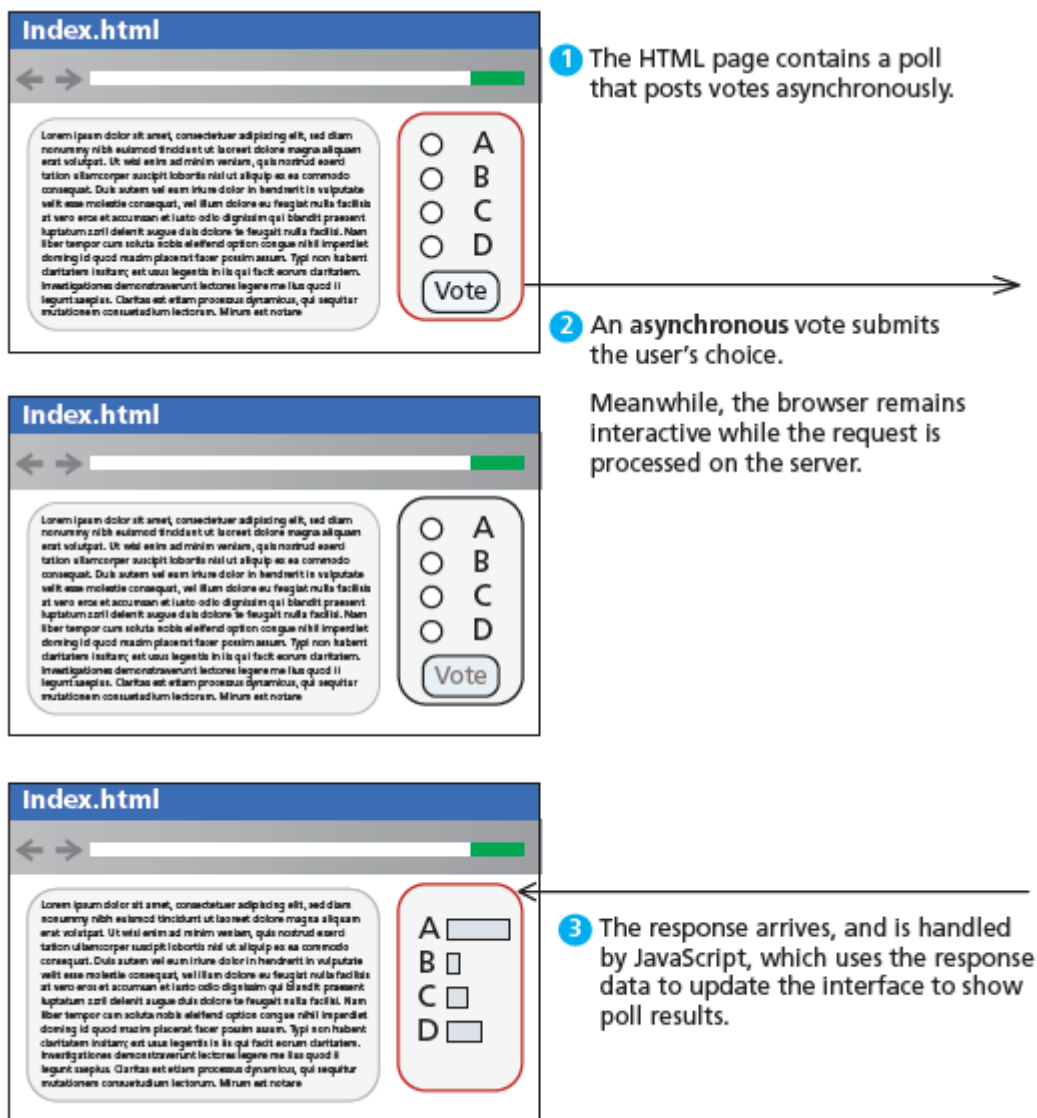


FIGURE 15.11 Illustration of a simple asynchronous web poll

Making a request to vote for option C in a poll could easily be encoded as a URL request `GET /vote.php?option=C`. However, rather than submit the whole page just to vote in the poll, jQuery's `$.get()` method sends that GET request asynchronously as follows:

```
$.get("/vote.php?option=C");
```

Note that the `$` symbol is followed by a dot. Recall that since `$` is actually shorthand for `jQuery()`, the above method call is equivalent to `jQuery().get("/vote.php?option=C");`

Attaching that function call to the form's submit event allows the form's default behavior to be replaced with an asynchronous GET request. `get()` method can request a resource very easily, handling the response from the request requires that we revisit the notion of the handler and listener. The event handlers used in jQuery are no different than those we've seen in JavaScript, except that they are attached to the event triggered by a request completing rather than a mouse move or key press. The formal definition of the `get()` method lists one required parameter `url` and three optional ones: `data`, a callback to a `success()` method, and a `dataType`.

```
jQuery.get ( url [, data ] [, success(data, textStatus, jqXHR) ] [, dataType ] )
```

- `url` is a string that holds the location to send the request.

- `data` is an optional parameter that is a query string or a *Plain Object*.

- `success(data, textStatus, jqXHR)` is an optional *callback* function that

executes when the response is received. Callbacks are the programming term given to placeholders for functions so that a function can be passed into another function and then called from there (called back). This callback function can take three optional parameters

- `data` holding the body of the response as a string.

- `textStatus` holding the status of the request (i.e., "success").

◦ jqXHR holding a jqXHR object, described shortly.

- `DataType` is an optional parameter to hold the type of data expected from the server. By default jQuery makes an intelligent guess between **xml**, **json**, **script**, or **html**.

```
$.get("/vote.php?option=C", function(data, textStatus, jsXHR) {  
  if (textStatus=="success") {  
    console.log("success! response is:" + data);  
  }  
  else {  
    console.log("There was an error code"+jsXHR.status);  
  }  
  console.log("all done");  
});
```

The callback function is passed as the second parameter to the `get()` method and uses the `textStatus` parameter to distinguish between a successful post and an error. The `data` parameter contains plain text and is echoed out to the user in an alert. Passing a function as a parameter can be an odd syntax for newcomers to jQuery. Unfortunately, if the page requested (`vote.php`, in this case) does not exist on the server, then the callback function does not execute at all, so the code announcing an error will never be reached. To address this we can make use of the `jqXHR` object to build a more complete solution.

The jqXHR Object

All of the `$.get()` requests made by jQuery return a **jqXHR** object to encapsulate the response from the server. In practice that means the `data` being referred to in the callback from Listing 15.13 is actually an object with backward compatibility with `XMLHttpRequest`. The following properties and methods are provided to conform to the `XMLHttpRequest` definition.

- `abort()` stops execution and prevents any callback or handlers from receiving the trigger to execute.
- `getResponseHeader()` takes a parameter and gets the current value of that header. `readyState` is an integer from 1 to 4 representing the state of the request. The values include 1: sending, 3: response being processed, and 4: completed.
- `responseXML` and/or `responseText` the main response to the request.
- `setRequestHeader(name, value)` when used before actually instantiating the request allows headers to be changed for the request.
- `status` is the HTTP request status codes described back in Chapter 1. (200 = ok)
- `statusText` is the associated description of the status code.

`jqXHR` objects have methods, `done()`, `fail()`, and `always()`, which allow us to structure our code in a more modular way than the inline callback.

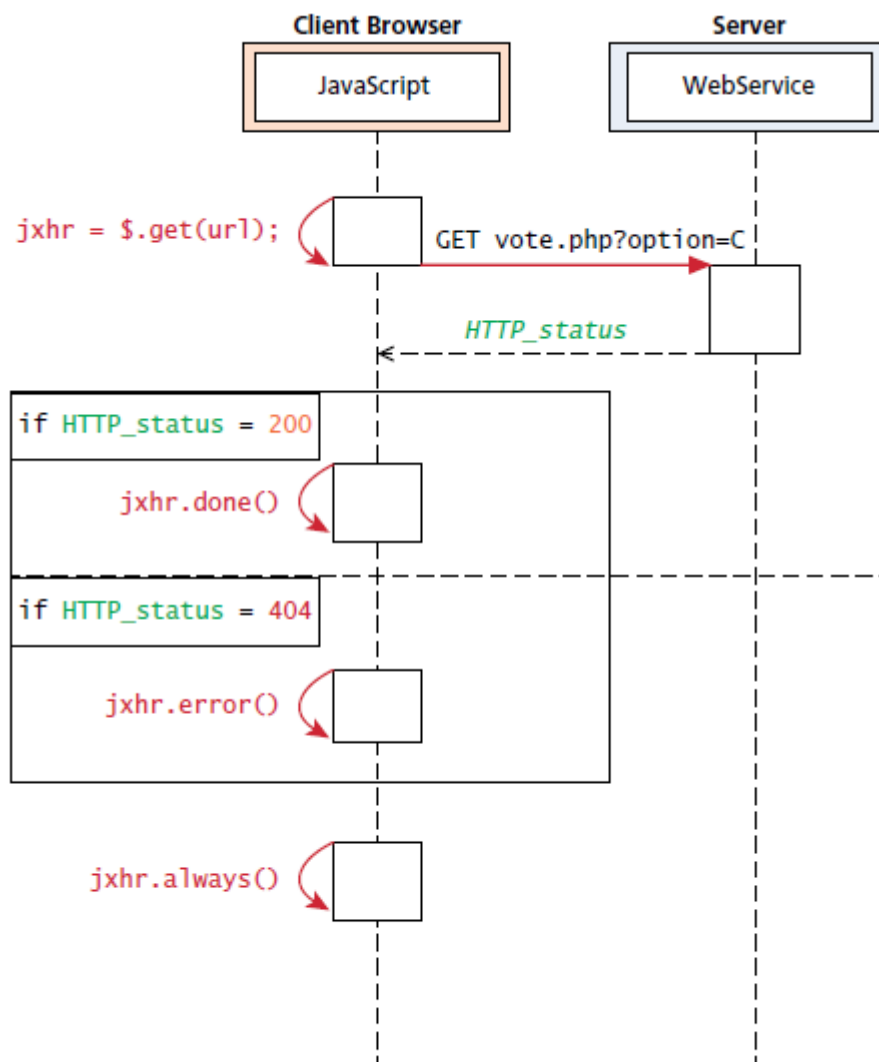


FIGURE 15.12 Sequence diagram depicting how the jqXHR object reacts to different response codes

POST Requests

POST requests are often preferred to GET requests because one can post an unlimited amount of data, and because they do not generate viewable URLs for each action. GET requests are typically not used when we have forms because of the messy URLs and that limitation on how much data we can transmit. Finally, with POST it is possible to transmit files, something which is not possible with GET. Although the differences between a GET and POST request are relatively minor, the HTTP 1.1 definition describes GET as a **safe method** meaning that they should not change anything, and should only read data. POSTs on the other hand are not safe, and should be used whenever we are changing the state of our system.

jQuery handles POST almost as easily as GET, with the need for an added field to hold our data. The formal definition of a jQuery `post()` request is identical to the `get()` request, aside from the method name.

jQuery.post (url [, data] [, success(data, textStatus, jqXHR)][, dataType])

The main difference between a POST and a GET http request is where the data is transmitted. The data parameter, if present in the function call, will be put into the body of the request. Interestingly, it can be passed as a string (with each name=value pair separated with a “&” character) like a GET request or as a Plain Object, as with the `get()` method.

If we were to convert our vote casting code from Listing 15.14 to a POST request, it would simply change the first line from

```
var jqxhr = $.get("/vote.php?option=C");
```

to

```
var jqxhr = $.post("/vote.php", "option=C");
```

Since jQuery can be used to submit a form, you may be interested in the shortcut method `serialize()`, which can be called on any form object to return its current key-value pairing as an & separated string, suitable for use with `post()`. Consider our simple vote-casting example. Since the poll’s form has a single field, it’s easy to understand the ease of creating a short query string on the fly. However, as forms increase in size this becomes more difficult, which is why jQuery includes a

helper function to serialize an entire form in one step. The `serialize()` method can be called on a DOM form element as follows:

```
var postData = $("#voteForm").serialize();
```

With the form's data now encoded into a query string (in the `postData` variable), you can transmit that data through an asynchronous POST using the `$.post()` method as follows:

```
$.post("vote.php", postData);
```

Complete Control over AJAX

It turns out both the `$.get()` and `$.post()` methods are actually shorthand forms for the `jQuery().ajax()` method, which allows fine-grained control over HTTP JavaScript requests including the modification of headers and use of cache controls. The `ajax()` method has two versions. In the first it takes two parameters: a URL and a Plain Object (also known as an object literal), containing any of over 30 fields. A second version with only one parameter is more commonly used, where the URL is but one of the key-value pairs in the Plain Object. The one line required to post our form using `get()` becomes the more verbose code.

```
$.ajax({ url: "vote.php",
  data: $("#voteForm").serialize(),
  async: true,
  type: post
});
```

To pass HTTP headers to the `ajax()` method, you enclose as many as you would like in a Plain Object. To illustrate how you could override User-Agent and Referer headers in the POST.

```
$.ajax({ url: "vote.php",
  data: $("#voteForm").serialize(),
  async: true,
  type: post,
  headers: {"User-Agent" : "Homebrew JavaScript Vote Engine agent",
    "Referer": "http://funwebdev.com"
  }
});
```

Cross-Origin Resource Sharing (CORS)

Cross-origin resource sharing (CORS) uses new headers in the HTML5 standard implemented in most new browsers. If a site wants to allow any domain to access its content through JavaScript, it would add the following header to all of its responses. **Access-Control-Allow-Origin: ***

The browser, seeing the header, permits any cross-origin request to proceed (since `*` is a wildcard) thus allowing requests that would be denied otherwise (by default).

A better usage is to specify specific domains that are allowed, rather than cast the gates open to each and every domain. In our example the more precise header

```
Access-Control-Allow-Origin: www.funwebdev.com
```

will prevent all cross-site requests, except those originating from www.funwebdev.com, allowing content to be shared between domains as needed.

4.10 jQuery Foundations

A **library** or **framework** is software that you can utilize in your own software, which provides some common implementations of standard ideas. A web framework can be expected to have features related to the web including HTTP headers, AJAX, authentication, DOM manipulation, cross-browser implementations, and more.

Including jQuery in Your Page

Since the entire library exists as a source JavaScript file, importing jQuery for use in your application is as easy as including a link to a file in the `<head>` section of your HTML page. You must either link to a locally hosted version of the library or use an approved third-party host, such as Google, Microsoft, or jQuery itself.

Using a third-party **content delivery network (CDN)** is advantageous for several reasons. Firstly, the bandwidth of the file is offloaded to reduce the demand on your servers. Secondly, the user may already have cached the third-party file and thus not have to download it again, thereby reducing the total loading time. This probability is increased when using a CDN like Google rather than a developer-focused CDN like jQuery.

A disadvantage to the third-party CDN is that your jQuery will fail if the third party host fails, although that is unlikely given the mission-critical demands of large companies like Google and Microsoft.

To achieve the benefits of the CDN and increase reliability on the rare occasion it might be down, you can write a small piece of code to check if the first attempt to load jQuery was successful. If not, you can load the locally hosted version. This setup should be included in the `<head>` section of your HTML page as shown in Listing 15.5.

```
<script src="http://code.jquery.com/jquery-1.9.1.min.js"></script>
<script type="text/javascript">
window.jQuery ||
document.write('<script src="/jquery-1.9.1.min.js"></script>');
</script>
```

jQuery Selectors

Basic Selectors

The four basic selectors were defined, and include the universal selector, class selectors, id selectors, and elements selectors. To review:

- `$("*")` **Universal selector** matches all elements (and is slow).
- `$("tag")` **Element selector** matches all elements with the given element name.
- `$(".class")` **Class selector** matches all elements with the given CSS class.
- `$("#id")` **Id selector** matches all elements with a given HTML id attribute.

For example, to select the single `<div>` element with `id="grab"` you would write:

```
var singleElement = $("#grab");
```

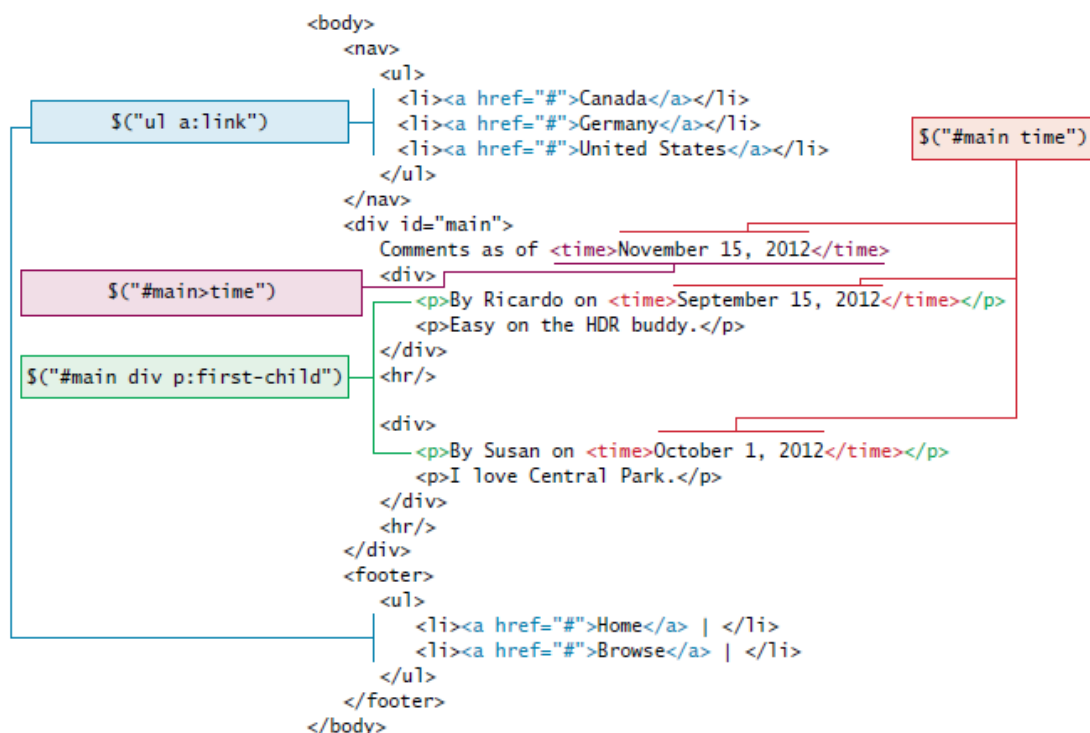
To get a set of all the `<a>` elements the selector would be:

```
var allAs = $("a");
```

These selectors are powerful enough that they can replace the use of `getElementById()` entirely.

Attribute Selector

An **attribute selector** provides a way to select elements by either the presence of an element attribute or by the value of an attribute.



A list of sample CSS attribute selectors was given in Chapter 3 (Table 3.4), but to jog your memory with an example, consider a selector to grab all `` elements with an `src` attribute beginning with

```
var artistImages = $('img[src^="/artist/"]');
```

([attribute⁼value], [attribute\$=value], [attribute*=value]).

Pseudo-elements are special elements, which are special cases of regular ones. these **pseudo-element selectors** allow you to append to any selector using the colon and one of :link, :visited, :focus, :hover, :active, :checked, :first-child, :first-line, and :first-letter.

```
var visitedLinks = $("a:visited");
```

Another powerful CSS selector included in jQuery's selection mechanism is the **contextual selectors** introduced in Chapter 3. These selectors allowed you to specify elements with certain relationships to one another in your CSS. These relationships included descendant (space), child (>), adjacent sibling (+), and general sibling (~). To select all <p>elements inside of <div>elements you would write

Content Filters

```
var allWarningText = $('body *:contains('warning')');
```

It will return a list of all the DOM elements with the word *warning* inside of them. You might imagine how we may want to highlight those DOM elements by coloring the background red as shown in Figure 15.5 with one line of code:

```
$("body *:contains('warning')).css("background-color", "#aa0000");
```

Selector	CSS Equivalent	Description
<code>\$(:button)</code>	<code>\$("button, input[type='button']")</code>	Selects all <i>buttons</i> .
<code>\$(:checkbox)</code>	<code>\$("[type=checkbox]")</code>	Selects all <i>checkboxes</i> .
<code>\$(:checked)</code>	No equivalent	Selects elements that are checked. This includes radio buttons and checkboxes.
<code>\$(:disabled)</code>	No equivalent	Selects form elements that are disabled. These could include <code><button></code> , <code><input></code> , <code><optgroup></code> , <code><option></code> , <code><select></code> , and <code><textarea></code>
<code>\$(:enabled)</code>	No equivalent	Opposite of <code>:disabled</code> . It returns all elements where the disabled attribute=false as well as form elements with no disabled attribute.
<code>\$(:file)</code>	<code>\$("[type=file]")</code>	Selects all elements of type file.
<code>\$(:focus)</code>	<code>\$(document.activeElement)</code>	The element with focus.
<code>\$(:image)</code>	<code>\$("[type=image]")</code>	Selects all elements of type image.
<code>\$(:input)</code>	No equivalent	Selects all <code><input></code> , <code><textarea></code> , <code><select></code> , and <code><button></code> elements.
<code>\$(:password)</code>	<code>\$("[type=password]")</code>	Selects all password fields.
<code>\$(:radio)</code>	<code>\$("[type=radio]")</code>	Selects all radio elements.
<code>\$(:reset)</code>	<code>\$("[type=reset]")</code>	Selects all the reset buttons.
<code>\$(:selected)</code>	No equivalent	Selects all the elements that are currently selected of type <code><option></code> . It does not include checkboxes or radio buttons.
<code>\$(:submit)</code>	<code>\$("[type=submit]")</code>	Selects all submit input elements.
<code>\$(:text)</code>	No equivalent	Selects all input elements of type text. <code>\$("[type=text]")</code> is almost the same, except that <code>\$(:text)</code> includes <code><input></code> fields with no type specified.

TABLE 15.1 jQuery form selectors and their CSS equivalents when applicable

Form Selectors

Since form HTML elements are well known and frequently used to collect and transmit data, there are jQuery selectors written especially for them. These selectors, listed in Table 15.1, allow for quick access to certain types of field as well as fields in certain states. attributes like the href attribute of an `<a>` tag, the src attribute of an ``, or the class attribute of most elements. In jQuery we can both set and get an attribute value by using the `attr()` method on any element from a selector. This function takes a parameter to specify which attribute, and the optional second parameter is the value to set it to. If no second parameter is passed, then the return value of the call is the current value of the attribute. Some example usages are:

```
// var link is assigned the href attribute of the first <a> tag
var link = $("a").attr("href");
// change all links in the page to http://funwebdev.com
$("a").attr("href", "http://funwebdev.com");
// change the class for all images on the page to fancy
$("img").attr("class", "fancy");
```

HTML Properties

Many HTML tags include properties as well as attributes, the most common being the *checked* property of a radio button or checkbox. In early versions of jQuery, HTML properties could be set using the `attr()` method. However, since properties are not technically attributes, this resulted in odd behavior. The `prop()` method is now the preferred way to retrieve and set the value of a property although, `attr()` may return some (less useful) values.

To illustrate this subtle difference, consider a DOM element defined by

```
<input class="meh" type="checkbox" checked="checked">
```

The value of the `attr()` and `prop()` functions on that element differ as shown below.


```
var theBox = $(".meh");
theBox.prop("checked") // evaluates to TRUE
theBox.attr("checked") // evaluates to "checked"
```

Changing CSS

Changing a CSS style is syntactically very similar to changing attributes. jQuery provides the extremely intuitive `css()` methods. There are two versions of this method (with two different method signatures), one to get the value and another to set it. The first version takes a single parameter containing the CSS attribute whose value you want and returns the current value.

```
$color = $("#colourBox").css("background-color"); // get the color
```

To modify a CSS attribute you use the second version of `css()`, which takes two parameters: the first being the CSS attribute, and the second the value.

```
// set color to red
```

```
$("#colourBox").css("background-color", "#FF0000");
```

Shortcut Methods

The `html()` method is used to get the HTML contents of an element (the part between the `<>` and `</>` tags associated with the `innerHTML` property in JavaScript). If passed with a parameter, it updates the HTML of that element. The `html()` method should be used with caution since the inner HTML of a DOM element can itself contain nested HTML elements! When replacing DOM with text, you may inadvertently introduce DOM errors since no validation is done on the new content (the browser wouldn't want to presume).

You can enforce the DOM by manipulating `TextNode` objects and adding them as children to an element in the DOM tree rather than use `html()`. While this enforces the DOM structure, it does complicate code. To illustrate, consider that you could replace the content of every `<p>` element with "jQuery is fun," with the one line of code:

```
$("p").html("jQuery is fun");
```

The shortcut methods `addClass(className)` / `removeClass(className)` add or remove a CSS class to the element being worked on. The `className` used for these functions can contain a space-separated list of class names to be added or removed. The `hasClass(className)` method returns true if the element has the `className` currently assigned. False, otherwise. The `toggleClass(className)` method will add or remove the class `className`, depending on whether it is currently present in the list of classes. The `val()` method returns the value of the element. This is typically used to retrieve values from input and select fields.

jQuery Listeners

Set Up after Page Load

In JavaScript, you learned why having your **listeners** set up inside of the `window.onload()` event was a good practice. Namely, it ensured the entire page and all DOM elements are loaded before trying to attach listeners to them. With jQuery we do the same thing but use the `$(document).ready()` event.

```
$(document).ready(function(){
//set up listeners on the change event for the file items.
$("input[type=file]").change(function(){
console.log("The file to upload is "+ this.value);
});
});
```

What is really happening is we are attaching our code to the **handler** for the `document.ready` event, which triggers when the page is fully downloaded and parsed into its DOM representation.

Listener Management

Setting up listeners for particular events is done in much the same way as JavaScript.

While pure JavaScript uses the `addEventListener()` method, jQuery has `on()` and `off()` methods as well as shortcut methods to attach events. Modifying the code in Listing 15.6 to use listeners rather than one handler yields the more modular code in below example: Note that the shortcut `:file` selector is used in place of the equivalent `input[type=file]`.

```

$(document).ready(function(){
$(":file").on("change",alertFileName); // add listener
});
// handler function using this
function alertFileName() {
console.log("The file selected is: "+this.value);
}

```

listing 15.7 Using the listener technique in jQuery with on and off methods

Modifying the DOM

jQuery comes with several useful methods to manipulate the DOM elements themselves. We have already seen how the `html()` function can be used to manipulate the inner contents of a DOM element and how `attr()` and `css()` methods can modify the internal attributes and styles of an existing DOM element.

Creating DOM and textNodes

If you decide to think about your page as a DOM object, then you will want to manipulate the tree structure rather than merely manipulate strings. Thankfully, jQuery is able to convert strings containing valid DOM syntax into DOM objects Automatically. creating a DOM node in JavaScript uses the `createElement()` method:

```
var element = document.createElement('div'); //create a new DOM node
```

However, since the jQuery methods to manipulate the DOM take an HTML string, jQuery objects, or DOM objects as parameters, you might prefer to define your element as

```
var element = $("

</div>"); //create new DOM node based on html


```

This way you can apply all the jQuery functions to the object, rather than rely on pure JavaScript, which has fewer shortcuts. If we consider creation of a simple `<a>` element with multiple attributes, you can see the comparison of the JavaScript and jQuery techniques in below example:

// pure JavaScript way

```

var jsLink = document.createElement("a");
jsLink.href = "http://www.funwebdev.com";
jsLink.innerHTML = "Visit Us";
jsLink.title = "JS";

```

// jQuery way

```

var jQueryLink = $("

```

// jQuery long-form way

```

var jQueryVerboseLink = $("

```

Prepending and Appending DOM Elements

When an element is defined in any of the ways described above, it must be inserted into the existing DOM tree. You can also insert the element into several places at once if you desire, since selectors can return an array of DOM elements. The `append()` method takes as a parameter an HTML string, a DOM object, or a jQuery object. That object is then added as the last child to the element(s) being selected. In Figure 15.6 we can see the effect of an `append()` method call. Each element with a class of `linkOut` has the `jsLink` element defined in below example appended to it.

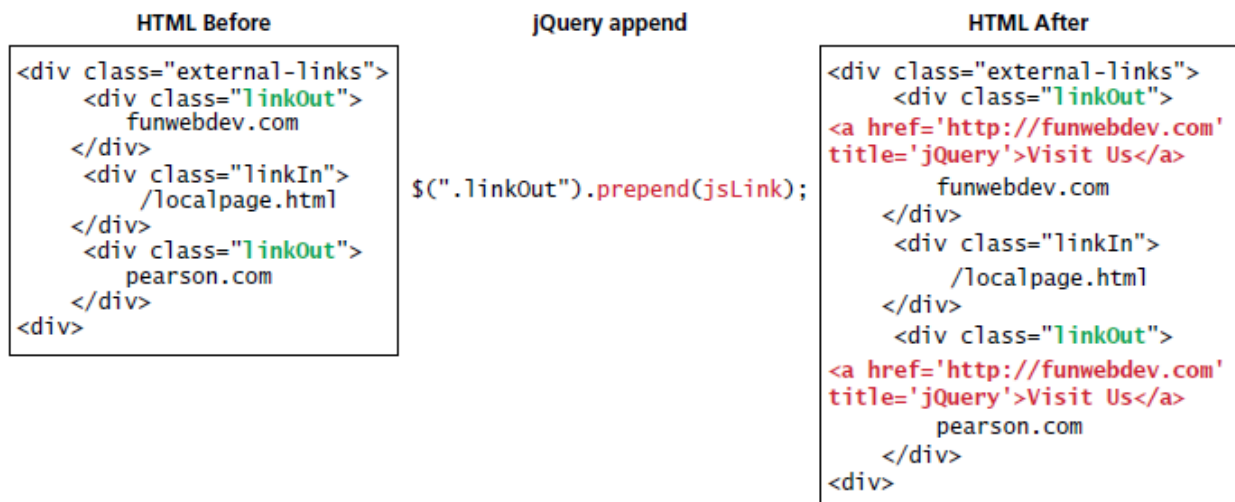


FIGURE 15.7 Illustration of `prepend()` adding a `` node

The `appendTo()` method is similar to `append()` but is used in the syntactically converse way. If we were to use `appendTo()`, we would have to switch the object making the call and the parameter to have the same effect as the previous code:

```
jsLink.appendTo($(".linkOut"));
```

The `prepend()` and `prependTo()` methods operate in a similar manner except that they add the new element as the first child rather than the last. See Figure 15.7 for an illustration of what happens with `prepend()`.

Wrapping Existing DOM in New Tags

One of the most common ways you can enhance a website that supports JavaScript is to add new HTML tags as needed to support some jQuery functions. Imagine for illustration purposes our art galleries being listed alongside some external links as described by the HTML in below example:

```

<div class="external-links">
  <div class="gallery">Uffuzi Museum</div>
  <div class="gallery">National Gallery</div>
  <div class="link-out">funwebdev.com</div>
</div>

```

If we wanted to wrap all the gallery items in the whole page inside, another `<div>` (perhaps because we wish to programmatically manipulate these items later) with class `galleryLink` we could write:

```
$(".gallery").wrap('<div class="galleryLink"/>');
```

which modifies the HTML to that shown in Listing 15.10. Note how each and every link is wrapped in the correct opening and closing and uses the `galleryLink` class.

```

<div class="external-links">
  <div class="galleryLink">
    <div class="gallery">Uffuzi Museum</div>
  </div>
  <div class="galleryLink">
    <div class="gallery">National Gallery</div>
  </div>
  <div class="link-out">funwebdev.com</div>
</div>

```

In a related demonstration of how succinctly jQuery can manipulate HTML, consider the situation where you wanted to add a title element to each `<div>` element that reflected the unique contents inside. To achieve this more sophisticated manipulation, you must pass a function as a parameter rather than a tag to the `wrap()` method, and that function will return a dynamically created `<div>` element as shown in below example:

```

$(".contact").wrap(function(){
  return "<div class='galleryLink' title='Visit ' + $(this).html() +
  "'></div>";
});

```

The `wrap()` method is a callback function, which is called for each element in aset (often an array). Each element then becomes this for the duration of one of the `wrap()` function's executions, allowing the unique title attributes as shown in below example:

```
<div class="external-links">
  <div class="galleryLink" title="Visit Uffuzi Museum">
    <div class="gallery">Uffuzi Museum</div>
  </div>
  <div class="galleryLink" title="Visit National Gallery">
    <div class="gallery">National Gallery</div>
  </div>
  <div class="link-out">funwebdev.com</div>
</div>
```

As with almost everything in jQuery, there is an inverse method to accomplish the opposite task. In this case, `unwrap()` is a method that does not take any parameters and whereas `wrap()` *added* a parent to the selected element(s), `unwrap()` *removes* the selected item's parent. Other methods such as `wrapAll()` and `wrapInner()` provide additional controls over wrapping DOM elements. The details of those methods can be found in the online jQuery documentation.

4.11 Animation

Animation Shortcuts

Animation is no different with a raw `animate()` method and many more easy-to-use shortcuts like `fadeIn()/fadeOut()`, `slideUp()/slideDown()`. One of the common things done in a dynamic web page is to show and hide an element. Modifying the visibility of an element can be done using `css()`, but that causes an element to change instantaneously, which can be visually jarring. To provide a more natural transition from hiding to showing, the `hide()` and `show()` methods allow developers to easily hide elements gradually, rather than through an immediate change.

The `hide()` and `show()` methods can be called with no arguments to perform a default animation. Another version allows two parameters: the duration of the animation (in milliseconds) and a callback method to execute on completion.

```
<div class="contact">
  <p>Randy Connolly</p>
  <div class="email">Show email</div>
</div>
<div class="contact">
  <p>Ricardo Hoar</p>
  <div class="email">Show email</div>
</div>
<script type='text/javascript'>
$( ".email" ).click(function() {
// Build email from 1st letter of first name + lastname
// @ mtroyal.ca
var fullName = $(this).prev().html();
var firstName = fullName.split(" ")[0];

var address = firstName.charAt(0) + fullName.split(" ")[1] +

"@mtroyal.ca";
$(this).hide(); // hide the clicked icon.
$(this).html("<a href='mailto:'"+address+"'>Mail Us</a>");
$(this).show(1000); // slowly show the email address.
});
</script>
```

listing 15.22 jQuery code to build an email link based on page content and animate its appearance `fadeIn()/fadeOut()`

The `fadeIn()` and `fadeOut()` shortcut methods control the opacity of an element. The parameters passed are the duration and the callback, just like `hide()` and `show()`. Unlike `hide()` and `show()`, there is no scaling of the element, just strictly control over the transparency.

`slideDown()/slideUp()`

The final shortcut methods we will talk about are `slideUp()` and `slideDown()`. These methods do not touch the opacity of an element, but rather gradually change its height.

Show email



Figure 15.15 Illustration of the show() animation using the icon from openiconlibrary.
sourceforge.net

Show email



Figure 15.16 Illustration of a fadeIn() animation



Figure 15.17 Illustration of the slideDown() animation

Toggle Methods

To toggle between the visible and hidden states (i.e., between using the hide() and show() methods), you can use the toggle() methods. To toggle between fading in and fading out, use the fadeToggle() method; toggling between the two sliding states can be achieved using the slideToggle() method. Using a toggle method means you don't have to check the current state and then conditionally call one of the two methods; the toggle methods handle those aspects of the logic for you.

Raw Animation

\$.get() and \$.post() methods are shortcuts for the complete \$.ajax() method, the animations shown this far are all specific versions of the generic animate() method. When you want to do animation that differs from the prepackaged animations, you will need to make use of animate. The animate() method has several versions, but the one we will look at has the following form:

.animate(properties, options);

The properties parameter contains a Plain Object with all the CSS styles of the final state of the animation. The options parameter contains another Plain Object with any of the options below set.

- always is the function to be called when the animation completes or stops with a fail condition. This function will always be called (hence the name).
- done is a function to be called when the animation completes.
- duration is a number controlling the duration of the animation.
- fail is the function called if the animation does not complete.
- progress is a function to be called after each step of the animation.
- queue is a Boolean value telling the animation whether to wait in the queue of animations or not. If false, the animation begins immediately.
- step is a function you can define that will be called periodically while the animation is still going. It takes two parameters: a now element, with the current numerical value of a CSS property, and an fx object, which is a temporary object with useful properties like the CSS attribute it represents (called tween in jQuery). See Listing 15.23 for example usage to do rotation.
- Advanced options called easing and special Easing allow for advanced control over the speed of animation. Movement rarely occurs in a linear fashion in nature. A ball thrown in the air slows down as it reaches the apex then accelerates toward the ground. In web development, easing functions are used to simulate that natural type of movement. They are mathematical equations that describe how fast or slow the transitions occur at various points during the animation. Included in jQuery are linear and swing easing functions. Linear is a straight line and so animation occurs at the same rate throughout while swing starts slowly and ends slowly. Figure 15.18 shows graphs for both the linear and swing easing functions. Easing functions are just mathematical definitions. For example, the function defining swing for values of time t between 0 and 1 is $\text{swing}(t) = - (1/2) \cos(\pi t) + 0.5$

The jQuery UI extension provides over 30 easing functions, including cubic functions and bouncing effects, so you should not have to define your own.

Show email



Figure 15.19 Illustration of rotation using the animate() method in Listing 15.23

It should be noted that step() callbacks are only made for CSS values that are numerical. This is why you often see a dummy CSS value used to control an unrelated CSS option like rotation

```
$(this).animate(  
  // parameter one: Plain Object with CSS options.  
  {opacity:"show","fontSize":"120%","marginRight":"100px"},  
  // parameter 2: Plain Object with other options including a  
  // step function  
  {step: function(now, fx) {  
    // if the method was called for the margin property  
    if (fx.prop=="marginRight") {  
      var angle=(now/100)*360; //percentage of a full circle  
      // Multiple rotation methods to work in multiple browsers  
      $(this).css("transform","rotate("+angle+"deg");  
      $(this).css("-webkit-transform","rotate("+angle+"deg");  
      $(this).css("-ms-transform","rotate("+angle+"deg");  
    }  
  },  
  duration:5000, "easing":"linear"  
});
```

listing 15.23 Use of animate() with a step function to do CSS3 rotation

4.12 Backbone MVC Frameworks

Getting Started with Backbone.js

Backbone is an MVC framework that further abstracts JavaScript with libraries intended to adhere more closely to the MVC model. In Backbone, you build your client scripts around the concept of models. These models are often related to rows in the site's database and can be loaded, updated, and eventually saved back to the database using a REST interface. You must download the source for these libraries to your server, and reference them just as we've done with jQuery. Remember that the underscore library is also required, so a basic inclusion will look like:

```
<script src="underscore-min.js"></script>  
<script src="backbone-min.js"></script>  
  
<form id="publishAlbums" method="post" action="publish.php">  
<h1>Publish Albums</h1>  
<ul id="albums">  
<!-- The albums will appear here -->  
</ul>  
<p id="totalAlbums">Count: <span>0</span></p>  
<input type="submit" id="publish" value="Publish" />  
</form>
```

listing 15.24 HTML for an album publishing interface

The MVC pattern in Backbone will use a Model object to represent the TravelAlbum, a Collection object to manage multiple albums, and a View to render the HTML for the model, and instantiate and render the entire application

Backbone Models

The term models can be a challenging one to apply, since authors of several frameworks and software engineering patterns already use the term.

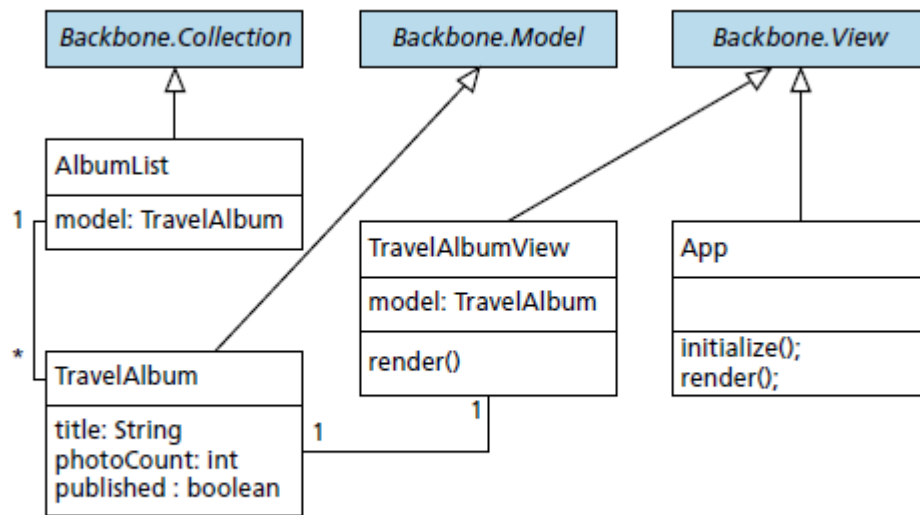


FIGURE 15.21 Illustration of Backbone Model, Collections, and Views for a Photo Album example

When using Backbone, you therefore begin by abstracting the elements you want to create models for. In our case, `TravelAlbum` will consist of a title, an image `photoCount`, and a Boolean value controlling whether it is published or not. The Models you define using Backbone must extend `Backbone.Model`, adding methods in the process as shown in Listing 15.25.

```
// Create a model for the albums
var TravelAlbum = Backbone.Model.extend({
  defaults: {
    title: 'NewAlbum',
    photoCount: 0,
    published: false
  },
  // Function to publish/unpublish
  toggle: function() {
    this.set('checked', !this.get('checked'));
  }
});
```

Collections

In addition to models, Backbone introduces the concept of Collections, which are normally used to contain lists of Model objects. These collections have advanced features and like a database can have indexes to improve search performance. In Listing 15.26, a collection of Albums, `AlbumList`, is defined by extending from Backbone's Collection object. In addition an initial list of `TravelAlbums`, named `albums`, is instantiated to illustrate the creation of some model objects inside a Collection.

```
// Create a collection of albums
var AlbumList = Backbone.Collection.extend({
  // Set the model type for objects in this Collection
  model: TravelAlbum,
  // Return an array only with the published albums
  GetChecked: function() {
    return this.where({checked: true});
  }
});
// Prefill the collection with some albums.
var albums = new AlbumList([
  new TravelAlbum({ title: 'Banff, Canada', photoCount: 42}),
  new TravelAlbum({ title: 'Santorini, Greece', photoCount: 102}),
]);
```

listing 15.26 Demonstration of a Backbone.js Collection defined to hold PhotoAlbums

Views

Views allow you to translate your models into the HTML that is seen by the users. They attach themselves to methods and properties of the Collection and define methods that will be called whenever Backbone determines the view needs refreshing.

For our example we extend a View as shown in Listing 15.27. In that code we attach our view to a particular tagName (in our case the element) and then var TravelAlbumView =

```
Backbone.View.extend({
  tagName: 'li',
  events: {
    'click': 'toggleAlbum'
  },
  initialize: function() {
    // Set up event listeners attached to change
    this.listenTo(this.model, 'change', this.render);
  },
  render: function() {
    // Create the HTML
    this.$el.html('<input type="checkbox" value="1" name="' +
      this.model.get('title') + '" /> ' +
      this.model.get('title') + '<span> ' +
      this.model.get('photoCount') + ' images</span>');
    this.$('input').prop('checked', this.model.get('checked'));
    // Returning the object is a good practice
    return this;
  },
  toggleAlbum: function() {
    this.model.toggle();
  }
});
```

listing 15.27 Deriving custom View objects for our model and Collection

associate the click event with a new method named toggleAlbum(). You must always override the render() method since it defines the HTML that is output. Finally, to make this code work you must also override the render of the main application. In our case we will base it initially on the entire <body> tag, and output our content based entirely on the models in our collection as shown in Listing 15.28.

4.13 HTML5 Web Storage

Web storage is a new JavaScript-only API introduced in HTML5.4 It is meant to be a replacement (or perhaps supplement) to cookies, in that web storage is managed by the browser; unlike cookies, web storage data is not transported to and from the server with every request and response. In addition, web storage is not limited to the 4K size barrier of cookies; the W3C recommends a limit of 5MB but browsers are allowed to store more per domain. Currently web storage is supported by current versions of the major browsers, including IE8 and above. However, since JavaScript, like cookies, can be disabled on a user's browser, web storage should not be used for mission-critical application functions.

Just as there were two types of cookies, there are two types of global web storage objects: localStorage and sessionStorage. The localStorage object is for saving information that will persist between browser sessions. The sessionStorage object is for information that will be lost once the browser session is finished.

These two objects are essentially key-value collections with the same interface (i.e., the same JavaScript properties and functions).

Using Web Storage

Below example illustrates the JavaScript code for writing information to web storage. Do note that it is not PHP code that interacts with the web storage mechanism but JavaScript. As demonstrated in the listing, there are two ways to store values in web storage: using the setItem() function, or using the property shortcut (e.g., sessionStorage.FavoriteArtist).

```

<form ... >
  <h1>Web Storage Writer</h1>
  <script language="javascript" type="text/javascript">

    if (typeof (localStorage) === "undefined" ||
        typeof (sessionStorage) === "undefined") {
      alert("Web Storage is not supported on this browser...");
    }
    else {
      sessionStorage.setItem("TodaysDate", new Date());
      sessionStorage.FavoriteArtist = "Matisse";

      localStorage.UserName = "Ricardo";
      document.write("web storage modified");
    }
  </script>
  <p><a href="WebStorageReader.php">Go to web storage reader</a></p>
</form>

```

Below example demonstrates that the process of reading from web storage is equally straightforward. The difference between sessionStorage and localStorage in this example is that if you close the browser after writing and then run the code in Below example, only the localStorage item will still contain a value.

```

<form id="form1" runat="server">
  <h1>Web Storage Reader</h1>
  <script language="javascript" type="text/javascript">

    if (typeof (localStorage) === "undefined" ||
        typeof (sessionStorage) === "undefined") {
      alert("Web Storage is not supported on this browser...");
    }
    else {
      var today = sessionStorage.getItem("TodaysDate");
      var artist = sessionStorage.FavoriteArtist;

      var user = localStorage.UserName;
      document.write("date saved=" + today);
      document.write("<br/>favorite artist=" + artist);
      document.write("<br/>user name = " + user);
    }
  </script>
</form>

```

Why Would We Use Web Storage?

Looking at the two previous listings you might wonder why we would want to use web storage. Cookies have the disadvantage of being limited in size, potentially disabled by the user, vulnerable to XSS and other security attacks, and being sent in every single request and response to and from a given domain. On the other hand, the fact that cookies are sent with every request and response is also their main advantage: namely, that it is easy to implement data sharing between the client browser and the server. Unfortunately with web storage, transporting the information within web storage back to the server is a relatively complicated affair involving the construction of a web service on the server (see Chapter 17) and then using asynchronous communication via JavaScript to push the information to the server.

A better way to think about web storage is not as a cookie replacement but as a local cache for relatively static items available to JavaScript. One practical use of web storage is to store static content downloaded asynchronously such as XML or JSON from a web service in web storage, thus reducing server load for subsequent requests by the session.

Figure 13.13 illustrates an example of how web storage could be used as a mechanism for reducing server data requests, thereby speeding up the display of the page on the browser, as well as reducing load on the server.

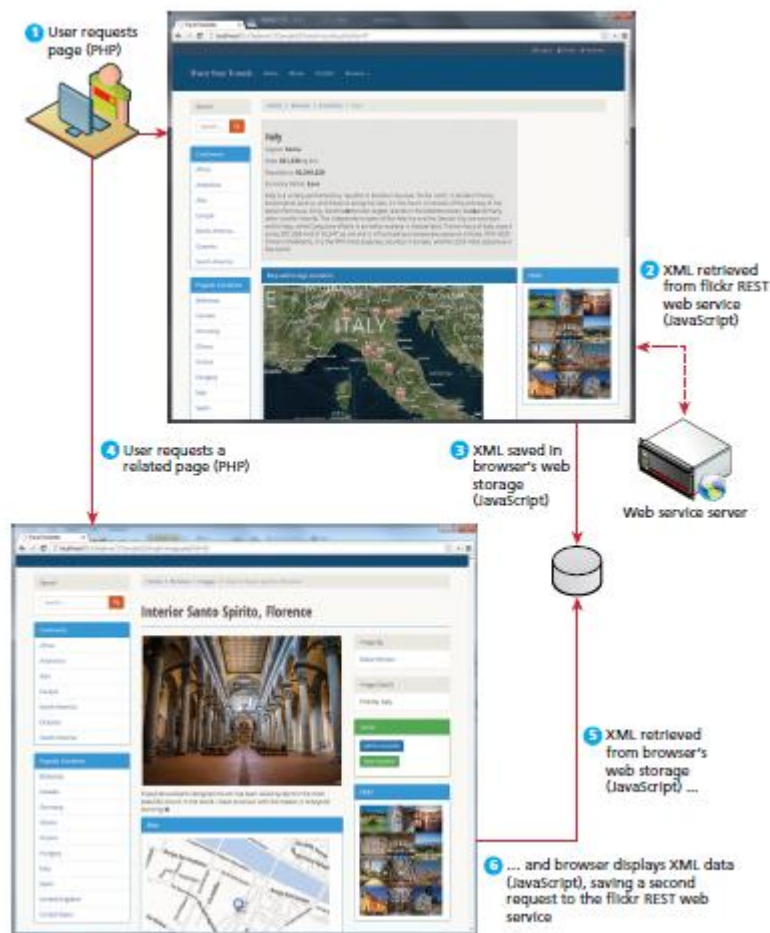


FIGURE 13.13 Using web storage

4.14 Caching

Caching is a vital way to improve the performance of web applications. Your browser uses caching to speed up the user experience by using locally stored versions of images and other files rather than re-requesting the files from the server. While important, from a server-side perspective, a server-side developer only has limited control over browser caching. There is a way, however, to integrate caching on the server side. Why is this necessary? Remember that every time a PHP page is requested, it must be fetched, parsed, and executed by the PHP engine, and the end result is HTML that is sent back to the requestor. For the typical PHP page, this might also involve numerous database queries and processing to build. If this page is being served thousands of times per second, the dynamic generation of that page may become unsustainable.

One way to address this problem is to cache the generated markup in server memory so that subsequent requests can be served from memory rather than from the execution of the page. There are two basic strategies to caching web applications. The first is page output caching, which saves the rendered output of a page or user control and reuses the output instead of reprocessing the page when a user requests the page again. The second is application data caching, which allows the developer to programmatically cache data.

Page Output Caching

In this type of caching, the contents of the rendered PHP page (or just parts of it) are written to disk for fast retrieval. This can be particularly helpful because it allows PHP to send a page response to a client without going through the entire page processing life cycle again (see Figure 13.14). Page output caching is especially useful for pages whose content does not change frequently but which require significant processing to create.

There are two models for page caching: **full page caching** and **partial page caching**. In full page caching, the entire contents of a page are cached. In partial page caching, only specific parts of a page are cached while the other parts are dynamically generated in the normal manner.

Page caching is not included in PHP by default, which has allowed a marketplace for free and commercial third-party cache add-ons such as Alternative PHP Cache (open source) and Zend (commercial) to flourish. However, one can easily create basic caching functionality simply by making use of the output buffering and time functions. The `mod_cache` module that comes with the Apache web server engine is the most common way websites implement page caching. This separates server tuning from your application code, simplifying development, and leaving cache control up to the web server rather than the application developer. It should be stressed that it makes no sense to apply page output caching to every page in a site. However, performance improvements can be gained (i.e., reducing server loads) by caching the page output of especially busy pages in which the content is the same for all users.

Application Data Caching

One of the biggest drawbacks with page output caching is that performance gains will only be had if the entire cached page is the same for numerous requests. However, many sites customize the content on each page for each user, so full or partial page caching may not always be possible.

An alternate strategy is to use application data caching in which a page will programmatically place commonly used collections of data that require time-intensive queries from the database or web server into cache memory, and then other pages that also need that same data can use the cache version rather than re-retrieve it from its original location. While the default installation of PHP does not come with an application caching ability, a widely available free PECL extension called memcache is widely used to provide this ability.⁵ Listing 13.10 illustrates a typical use of memcache. It should be stressed that memcache should not be used to store large collections. The size of the memory cache is limited, and if too many things are placed in it, its performance advantages will be lost as items get paged in and out. Instead, it should be used for relatively small collections of data that are frequently accessed on multiple pages.

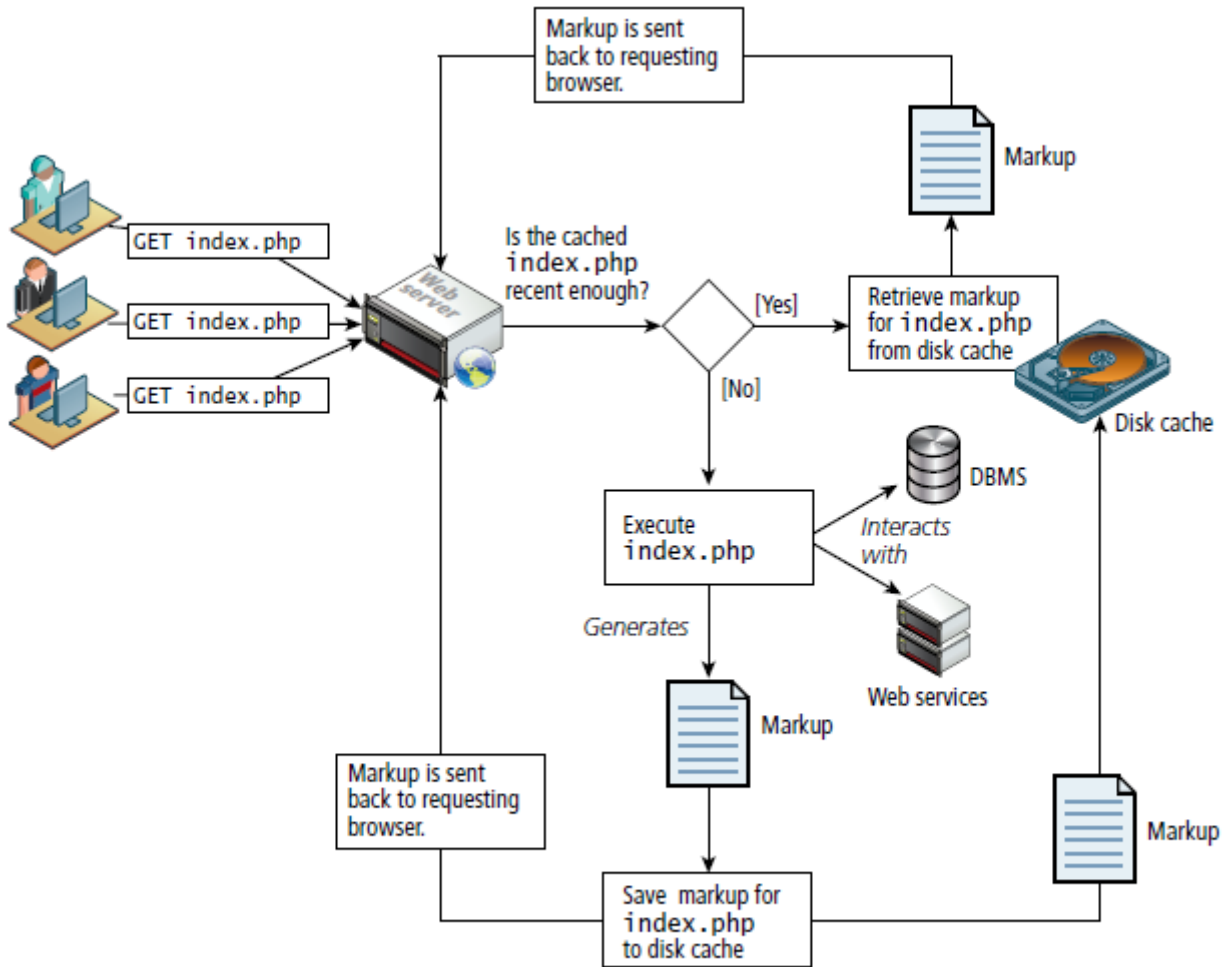


FIGURE 13.14 Page output caching

```

<?php

// create connection to memory cache
$memcache = new Memcache;
$memcache->connect('localhost', 11211)
    or die ("Could not connect to memcache server");

$cacheKey = 'topCountries';
/* If cached data exists retrieve it, otherwise generate and cache
it for next time */
if ( ! isset($countries = $memcache->get($cacheKey)) ) {

    // since every page displays list of top countries as links
    // we will cache the collection

    // first get collection from database
    $cgate = new CountryTableGateway($dbAdapter);
    $countries = cgate->getMostPopular();

    // now store data in the cache (data will expire in 240 seconds)
    $memcache->set($cacheKey, $countries, false, 240)
        or die ("Failed to save cache data at the server");
}
// now use the country collection
displayCountryList($countries);

?>

```

4.15 Asynchronous File Transmission

Asynchronous file transmission is one of the most powerful tools for modern web applications. In the days of old, transmitting a large file could require your user to wait idly by while the file uploaded, unable to do anything within the web interface. Since file upload speeds are almost always slower than download speeds, these transmissions can take minutes or even hours, destroying the feeling of a “real” application. Unfortunately jQuery alone does not permit asynchronous file uploads! However, using clever tricks and HTML5 additions, you too can use asynchronous file uploads.

For the following examples consider a simple file-uploading HTML form defined in below ex:

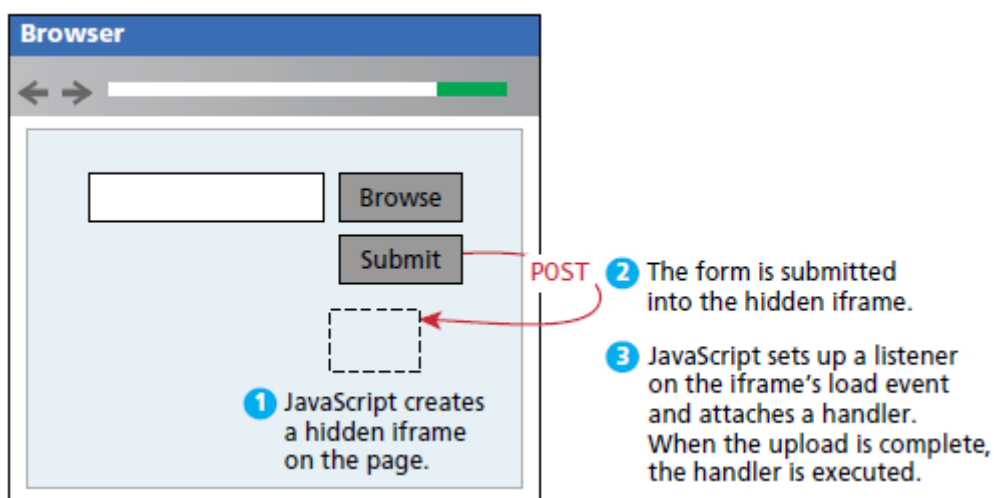
```

<form name="fileUpload" id="fileUpload" enctype="multipart/form-data"
    method="post" action="upload.php">
<input name="images" id="images" type="file" multiple />
<input type="submit" name="submit" value="Upload files!" />
</form>

```

Old iframe Workarounds

The original workaround to allow the asynchronous posting of files was to use a hidden <iframe> element to receive the posted files. Given that jQuery still does not natively support the asynchronous uploading of files, this technique persists to this day and may be found in older code you have to maintain. As illustrated in Figure



and Listing 15.18, a hidden `<iframe>` allows one to post synchronously to another URL in another window. If JavaScript is enabled, you can also hide the upload button and use the change event instead to trigger a file upload. You then use the `<iframe>` element's onload event to trigger an action when it is done loading. When the window is done loading, the file has been received and we use the return message to update our interface much like we do with AJAX normally.

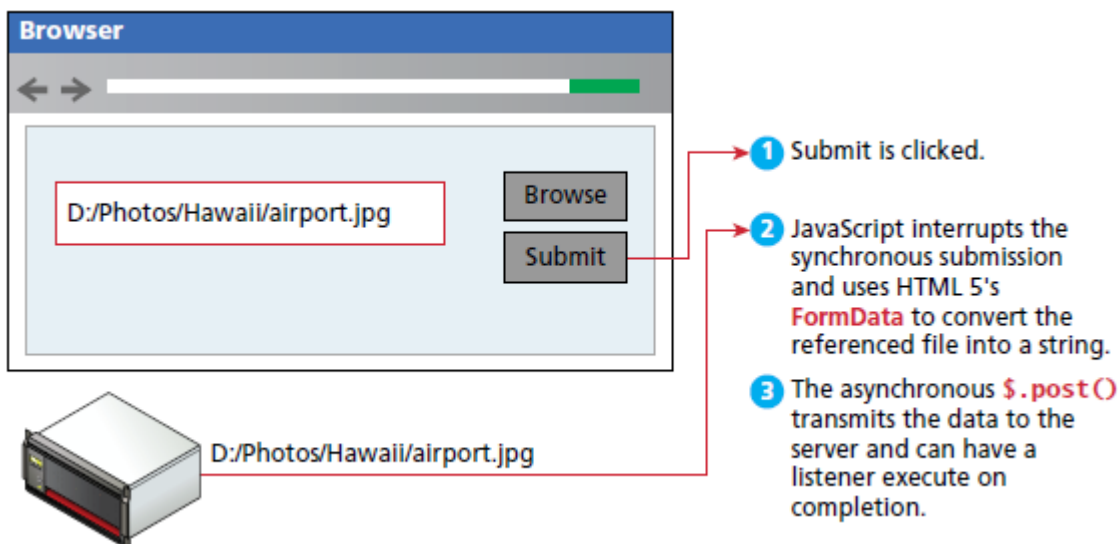
```
$(document).ready(function() {  
    // set up listener when the file changes  
    $(":file").on("change",uploadFile);  
    // hide the submit buttons  
    $("input[type=submit]").css("display","none");  
});  
  
// function called when the file being chosen changes  
function uploadFile () {  
    // create a hidden iframe  
    var hidName = "hiddenIFrame";  
    $("#fileUpload").append("<iframe id='"+hidName+"' name='"+hidName+"' style='display:none' src='#' ></iframe>");  
  
    // set form's target to iframe  
    $("#fileUpload").prop("target",hidName);  
    // submit the form, now that an image is in it.  
    $("#fileUpload").submit();  
}
```

This technique exploits the fact that browsers treat each `<iframe>` element as a separate window with its own thread. By forcing the post to be handled in another window, we don't lose control of our user interface while the file is uploading. Although it works, it's a workaround using the fact that every browser can post a file synchronously. A more modular and "pure" technique would be to somehow serialize the data in the file being uploaded with JavaScript and then post it in the body of a post request asynchronously.

Thankfully, the newly redefined XMLHttpRequest Level 2 (XHR2) specification allows us to get access to file data and more through the FormData interface

The FormData Interface

Using the FormData interface and File API, which is part of HTML5, you no longer have to trick the browser into posting your file data asynchronously. However, you are limited to modern browsers that implement the new specification.



The FormData interface provides a mechanism for JavaScript to read a file from the user's computer (once they choose the file) and encode it for upload. You can use this mechanism to upload a file asynchronously. Intuitively the browser is already able to do this, since it can access file data for transmission in synchronous posts. The FormData interface simply exposes this functionality to the developer, so you can turn a file into a string when you need to.

The `<iframe>` method `uploadFile()` from Listing 15.18 can be replaced with the more elegant and straightforward code in Listing 15.19. In this pure AJAX technique the form object is passed to a FormData constructor, which is then used in the call to `send()` the XHR2 object. This code attaches listeners for various events that may occur.

```
function uploadFile () {
    // get the file as a string
    var formData = new FormData($("#fileUpload")[0]);

    var xhr = new XMLHttpRequest();
    xhr.addEventListener("load", transferComplete, false);
    xhr.addEventListener("error", transferFailed, false);
    xhr.addEventListener("abort", transferCanceled, false);

    xhr.open('POST', 'upload.php', true);
    xhr.send(formData);           // actually send the form data

    function transferComplete(evt) { // stylized upload complete
        $("#progress").css("width", "100%");
        $("#progress").html("100%");
    }

    function transferFailed(evt) {
        alert("An error occurred while transferring the file.");
    }

    function transferCanceled(evt) {
        alert("The transfer has been canceled by the user.");
    }
}
```

Appending Files to a POST

When we consider uploading multiple files, you may want to upload a single file, rather than the entire form every time. To support that pattern, you can access a single file and post it by appending the raw file to a FormData object as shown in Listing 15.20. The advantage of this technique is that you submit each file to the server asynchronously as the user changes it; and it allows multiple files to be transmitted at once.

```
var xhr = new XMLHttpRequest();
// reference to the 1st file input field
var theFile = $(":file")[0].files[0];
var formData = new FormData();
formData.append('images', theFile);
```

It should be noted that back in Listing 15.17 the file input is marked as multiple, and so, if supported by the browser, the user can select many files to upload at once. To support uploading multiple files in our JavaScript code, we must loop through all the files rather than only hard-code the first one. Listing 15.21 shows a better script than Listing 15.20, since it handles multiple files being selected and uploaded at once.

```
var allFiles = $(":file")[0].files;
for (var i=0; i<allFiles.length; i++) {
    formData.append('images[]', allFiles[i]);
}
```

The main challenge of asynchronous file upload is that your implementation must consider the range of browsers being used by your users. While the new XHR2 specification and FormData interfaces are “pure” and easy to use, they are not widely supported yet across multiple platforms and browsers, making reliance on them bad practice. Conversely the <iframe> workaround works well on more browsers, but simply feels inelegant and perhaps not worthy of your support and investment of time.

4.16 Overview of Web Services

Web services are the most common example of a computing paradigm commonly referred to as service-oriented computing (SOC), which utilizes something called “services” as a key element in the development and operation of software applications.

A service is a piece of software with a platform-independent interface that can be dynamically located and invoked. Web services are a relatively standardized mechanism by which one software application can connect to and communicate with another software application using web protocols. Web services make use of the HTTP protocol so that they can be used by any computer with Internet connectivity. As well, web services typically use XML or JSON (which will be covered shortly) to encode data within HTTP transmissions so that almost any platform should be able to encode or retrieve the data contained within a web service. The benefit of web services is that they potentially provide interoperability between different software applications running on different platforms. Because web services use common and universally supported standards (HTTP and XML/JSON), they are supported on a wide variety of platforms. Another key benefit of web services is that they can be used to implement something called a service-oriented

architecture (SOA). This type of software architecture aims to achieve very loose coupling among interacting software services. The rationale behind an SOA is one that is familiar to computing practitioners with some experience in the enterprise: namely, how to best deal with the problem of application integration. Due to corporate mergers, longer-lived legacy applications, and the need to integrate with the Internet, getting different software applications to work together has become a major priority of IT organizations. SOA provides a very palatable potential solution to application integration issues. Because services are independent software entities, they can be offered by different systems within an organization as well as by different organizations. As such, web services can provide a computing infrastructure for application integration and collaboration within and between organizations, as shown in Figure 17.7.

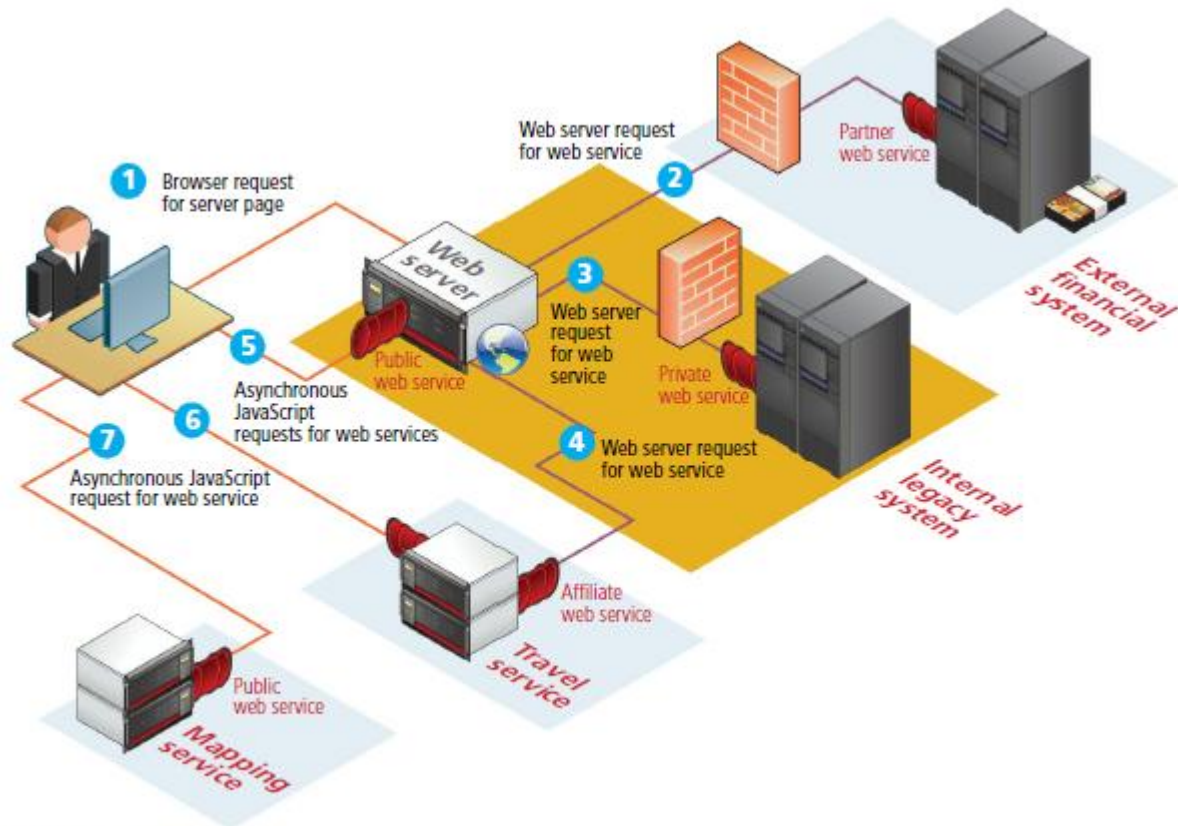


FIGURE 17.7 Overview of web services

SOAP Services

In the first iteration of web services fever, the attention was on a series of related XML vocabularies: WSDL, SOAP, and the so-called WS-protocol stack (WS-Security, WS-Addressing, etc.). In this model, WSDL is used to describe the operations and data types provided by the service. SOAP is the message protocol used to encode the service invocations and their return values via XML within the HTTP header, as can be seen in Figure 17.8. While SOAP and WSDL are complex XML schemas, this now relatively mature standard is well supported in the .NET and Java environments (perhaps a little less so with PHP). From the authors' professional and teaching experience, it is not necessary to have detailed knowledge of the SOAP and WSDL specifications to create and consume SOAP-based services. Using SOAP-based services is somewhat making to using a compiler: its output may be complicated to understand, but it certainly makes life easier for most programmers. Yet, despite the superb tool support in these two environments, by the middle years of the 2000s, the enthusiasm for SOAP-based web services had cooled.

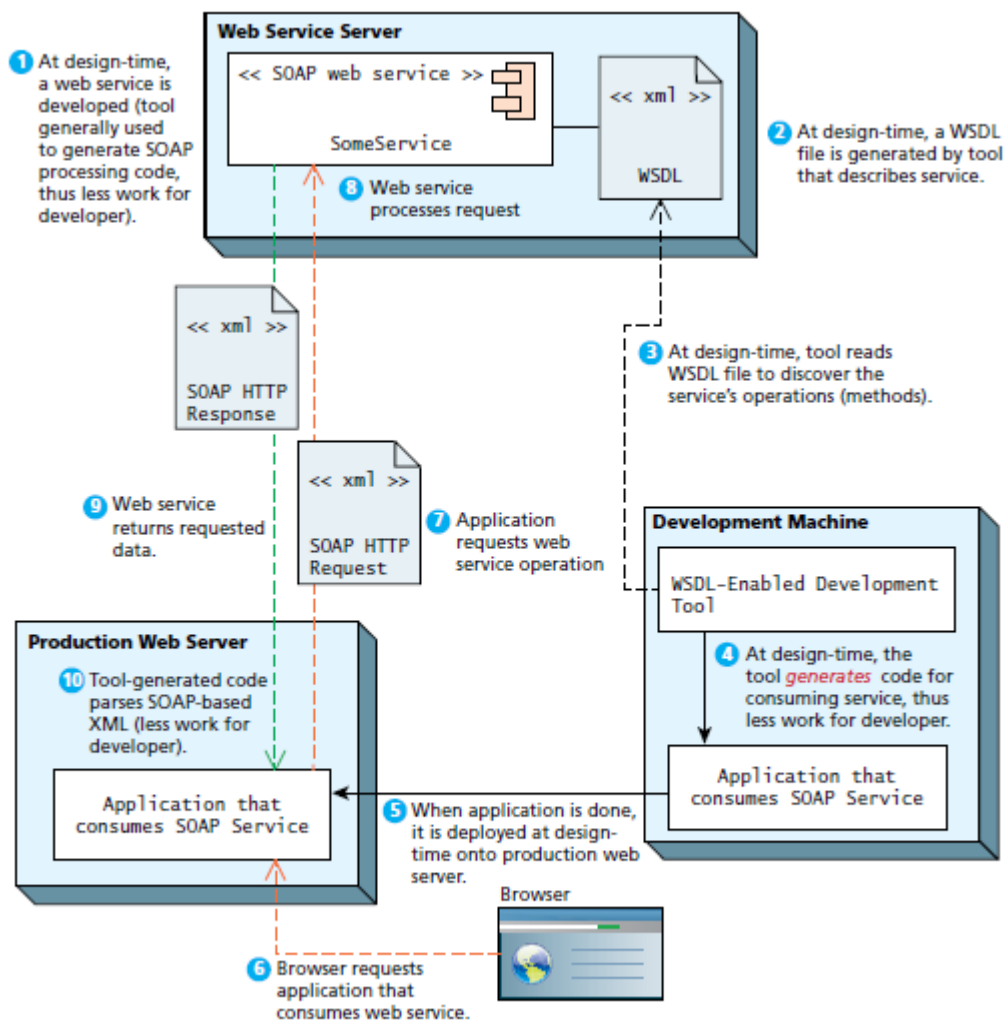


FIGURE 17.8 SOAP web services

REST Services

A RESTful web service does away with the service description layer as well as doing away with the need for a separate protocol for encoding message requests and responses. Instead it simply uses HTTP URLs for requesting a resource/object (and for encoding input parameters). The serialized representation of this object, usually an XML or JSON stream, is then returned to the requestor as a normal HTTP response. No special steps are needed to deploy a REST-based service, no special tools (other than a browser) are generally needed to test a RESTful service, and it is easier to scale for a large number of clients using well-established practices and experience with caching, clustering, and load balancing traditional dynamic HTTP websites.

With the broad interest in the asynchronous consumption of server data at the browser using JavaScript (generally referred to as AJAX) in the latter half of this decade, the lightweight nature of REST made it significantly easier to consume in JavaScript than SOAP. Indeed, if an object is serialized via JSON, it can be turned into a complex JavaScript object in one simple line of JavaScript. However, since many REST web services use XML as the data format, manual XML parsing and processing is required in order to deserialize a REST response back into a usable object, as shown in Figure 17.9. (With the SOAP approach, in contrast, tools can use the WSDL document to automatically generate proxy classes at development time, which in turn obviates the necessity of writing the serialize/deserialize code yourself.)

REST appears to have almost completely displaced SOAP services. For instance, in July 2013, the programmableweb.com API directory had 2030 indexed SOAP services in comparison to 6088 REST services. While some of the most popular services, such as those from Amazon, eBay, and Flickr, support both formats, others, such as Facebook, Google, YouTube, and Wikipedia, have either discontinued SOAP support or have never offered it. For this reason, this chapter will only cover the consumption and creation of REST-based services.

The relatively easy availability of a wide range of RESTful services has given rise to a new style of web development, often referred to as a mashup, which generally refers to a website that combines and integrates data from a variety of different sources. Even websites that are not overtly mashups nonetheless often make use of some external data via the consumption of REST services. The proliferation of maps, externalized search, Amazon widgets, and so on, on a wide variety of sites are examples of the commonality of the consumption of REST services.

An Example Web Service

Perhaps the best way to understand RESTful web service would be to examine a sample one. In this section we will look at the Google Geocoding API. The term geocoding typically refers to the process of turning a real-world address (such as British Museum, Great Russell Street, London, WC1B 3DG) into geographic coordinates, which are usually latitude and longitude values (such as 51.5179231, -0.1271022). Reverse geocoding is the process of converting geographic coordinates into a human-readable address.

The Google Geocoding API provides a way to perform geocoding operations via an HTTP GET request, and thus is an especially useful example of a RESTful web service.

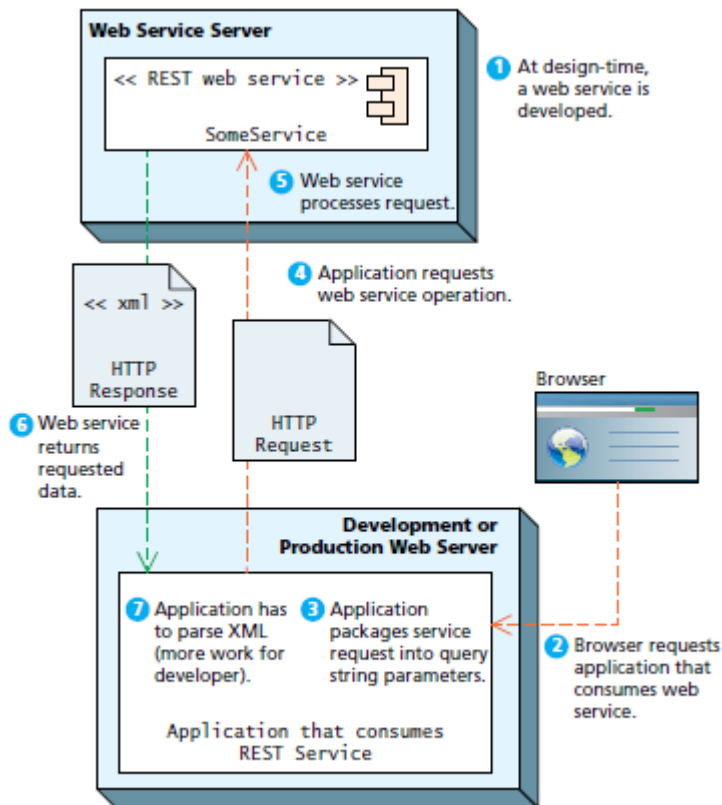


FIGURE 17.9 REST web services

Like all of the REST web services we will be examining in this chapter, using a web service begins with an HTTP request. In this case the request will take the following form:

<http://maps.googleapis.com/maps/api/geocode/xml?parameters>

The parameters in this case are address (for the real-world address to geocode) and sensor (for whether the request comes from a device with a location sensor).

So an example geocode request would look like the following:

<http://maps.googleapis.com/maps/api/geocode/xml?address=British%20Museum,+Great+Russell+Street,+London,+WC1B+3DG&sensor=false>

Notice that a REST request, like all HTTP requests, must URL encode special characters such as spaces. If the request is well formed and the service is working, it will return an HTTP response similar to that shown in Listing 17.13 (with some omissions and indenting spaces added for readability).


```

HTTP/1.1 200 OK
Content-Type: application/xml; charset=UTF-8
Date: Fri, 19 Jul 2013 19:15:54 GMT
Expires: Sat, 20 Jul 2013 19:15:54 GMT
Cache-Control: public, max-age=86400
Vary: Accept-Language
Content-Encoding: gzip
Server: mafe
Content-Length: 512
X-XSS-Protection: 1; mode=block
X-Frame-Options: SAMEORIGIN

<?xml version="1.0" encoding="UTF-8"?>
<GeocodeResponse>
  <status>OK</status>
  <result>
    <type>route</type>
    <formatted_address>
      Great Russell Street, London Borough of Camden, London, UK
    </formatted_address>
    <address_component>
      <long_name>Great Russell Street</long_name>
      <short_name>Great Russell St</short_name>
      <type>route</type>
    </address_component>
    <address_component>
      <long_name>London</long_name>
      <short_name>London</short_name>
      <type>locality</type>
      <type>political</type>
    </address_component>
    ...
    <geometry>
      <location>
        <lat>51.5179231</lat>
        <lng>-0.1271022</lng>
      </location>

      <location_type>GEOMETRIC_CENTER</location_type>
      ...
    </geometry>
  </result>
</GeocodeResponse>

```

After receiving this response, our program would then presumably need some type of XML processing in order to extract the latitude and longitude values (perhaps the simplest way to do so would be via XPath).

Identifying and Authenticating Service Requests

The previous section illustrated a sample request to a REST-based web service and its XML response. That particular service was openly available to any request (though its term of service license limited how the response data could be used). Most web services are not open in the same way. Instead, they typically employ one of the following techniques:

- Identity. Each web service request must identify who is making the request.
- Authentication. Each web service request must provide additional evidence that they are who they say they are.

Many web services are not providing information that is especially private or proprietary. For instance, the Flickr web service, which provides URLs to publicly available photos on their site in response to search criteria, is in some ways simply an XML version of the main site's already existing search facility. Since no private user data is being requested, it only expects each web service request to include one or more API keys to identify who is making the request.

This typically is done not only for internal record-keeping, but more importantly to keep service request volume at a manageable level. Most external web service APIs limit the number of web service requests that can be made, generally either per second, per hour, or per day. For instance, Panoramio limits requests to 100,000 per day while Google Maps and Microsoft Bing Maps allow 50,000 geocoding requests per day; Instagram allows 5000 requests per hour but Twitter allows just 100 to 400 requests per hour (it can vary); Amazon and last.fm limit requests to just one per second. Other services such as Flickr, NileGuide, and YouTube have no documented request limits.

Web services that make use of an API key typically require the user (i.e., the developer) to register online with the service for an API key. This API key is then added to the GET request as a query string parameter. For instance, a geocoding request to the Microsoft Bing Maps web service will look like the following (in this particular case, the actual Bing API key is a 64-character string):

```

http://dev.virtualearth.net/REST/v1/Locations?o=xml&query=British%20Museum,+Great+Russell+Street,+London,+WC1B+3DG,+UK&key=[BING API KEY HERE]

```

While some web services are simply providing information already available on their website, other web services are providing private/proprietary information or are involving financial transactions. In this case, these services not only may require an API key, but they also require some type of user name and password in order to perform an authorization.